

Inferring Graphics Programs from Images

Anonymous Author(s)

Affiliation

Address

email

Abstract

1

2 1 Introduction

3 How could an agent go from noisy, high-dimensional perceptual input to a symbolic, abstract object,
 4 like a computer program? Here we consider this problem within a graphics program synthesis domain.
 5 We develop an approach for converting natural images, such as hand drawings, into executable source
 6 code for drawing the original image. The graphics programs in our domain draw simple figures like
 7 those found in machine learning papers (see Fig.1). [The use of ‘graphics programs / visual programs’
 8 in the paper title, title of this section, and the body of this section feels too broad. ‘Graphics program’
 9 could conjur a lot of different ideas (esp. 3D graphics); don’t want to set the reader up to expect one
 10 thing and then be disappointed that what you’ve done isn’t that. You bring up diagram-drawing later
 11 in the intro; I think it should be made clear sooner (and certainly mentioned explicitly in the abstract,
 12 when you get around to writing that).]

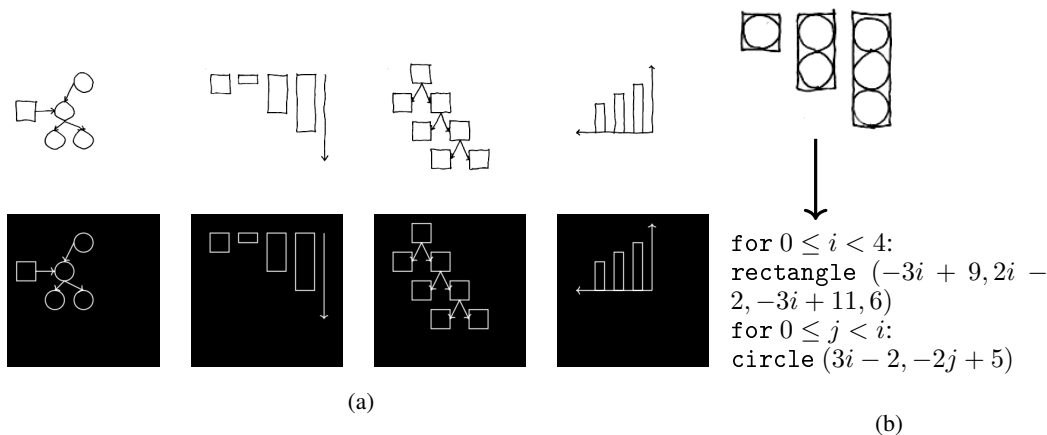


Figure 1: (a): Model learns to convert hand drawings (top) into LAT\textsubscript{E}X (bottom). (b) Synthesizes high-level *graphics program* from hand drawing.

13 High dimensional perceptual input may seem ill matched to the abstract semantics of a programming
 14 language. But programs with constructs like recursion or iteration produce a simpler *execution trace*
 15 of primitive actions; for our domain the primitive actions are drawing commands. Our hypothesis is
 16 that the execution trace of the program is better aligned with the perceptual input, and that the trace
 17 can act as a kind of bridge between perception and programs. We test this hypothesis by developing
 18 a model that learns to map from an image to the execution trace of the graphics program that drew
 19 it. With the execution trace in hand, we can bring to bear techniques from the program synthesis

community to recover the latent graphics program. This family of techniques, called *constraint-based program synthesis* [?], work by modeling a set of possible programs inside of a constraint solver, like a SAT or SMT solver [?]. These techniques excel at uncovering high-level symbolic structure, but are not well equipped to deal with real-valued perceptual inputs.

We develop a hybrid architecture for inferring graphics programs. Our approach uses a deep neural network infer an execution trace from an image; this network recovers primitive drawing operations such as lines, circles, or arrows, along with their parameters. For added robustness, we use the deep network as a proposal distribution for a stochastic search over execution traces. Finally, we use techniques in the program synthesis community to recover the program from its trace. The program synthesizer discovers constructs like loops and geometric operations like reflections and affine transformations. [This paragraph is all about making things a bit more specific, so you really need more specifics about program synth here.]

Each of these three components – the deep network, the stochastic search, the program synthesizer – confers its own advantages. From the deep network, we get a fast system that can recover plausible execution traces in about a minute [A minute seems slow to me, for deep net inference. Are you talking about training time, here, or...?]. From the stochastic search we get added robustness; essentially, the stochastic search can correct mistakes made by the deep network’s proposals. From the program synthesizer, we get abstraction: our system recovers coordinate transformations, for loops, and subroutines, which are useful for downstream tasks and can help correct some mistakes of the earlier stages. [I wonder if this would work even better as a bulleted list...]

2 Related work

Our work bears resemblance to the Attend-Infer-Repeat (AIR) system, which learns to decompose an image into its constituent objects [1]. AIR learns an iterative inference scheme which infers objects one by one and also decides when to stop inference; this is similar to our approach’s first stage, which parses images into program execution traces. Our approach further produces interpretable, symbolic programs which generate those execution traces. The two approaches also differ in their architectures and training regimes: AIR learns a recurrent auto-encoding model via variational inference, whereas our parsing stage learns an autoregressive-style model from randomly-generated (execution trace, image) pairs. Finally, while AIR was evaluated on multi-MNIST images and synthetic 3D scenes, we focus on parsing and interpreting hand-drawn sketches.

Our image-to-execution-trace parsing architecture builds on prior work on controlling procedural graphics programs [2]. Given a program which generates random 2D recursive structures such as vines, that system learns a structurally-identical “guide program” whose output can be directed, via neural networks, to resemble a given target image. We adapt this method to a different visual domain (figures composed of multiple objects), using a broad prior over possible scenes as the initial program and viewing the execution trace through the guide program as a symbolic parse of the target image. We then show how to synthesize higher-level programs from these execution traces.

In the computer graphics literature, there have been other systems which convert sketches into procedural representations. One uses a convolutional network to match a sketch to the output of a parametric 3D modeling system [4]. Another uses convolutional networks to support sketch-based instantiation of procedural primitives within an interactive architectural modeling system [5]. Both systems focus on inferring fixed-dimensional parameter vectors. In contrast, we seek to automatically learn a structured, programmatic representation of a sketch which captures higher-level visual patterns.

Prior work has also applied sketch-based program synthesis to authoring graphics programs. In particular, Sketch-n-Sketch presents a bi-directional editing system in which direct manipulations to a program’s output automatically propagate to the program source code [3]. We see this work as complementary to our own: programs produced by our method could be provided to a Sketch-n-Sketch-like system as a starting point for further editing.

[Do you also want to cite your own work on “Unsupervised Learning by Program Synthesis” here?]

69 3 Neural architecture for inferring drawing execution traces

70 We developed a deep network architecture for efficiently inferring a execution trace, T , from an
 71 image, I . Our model constructs the trace one drawing command at a time. When predicting the next
 72 drawing command, the network takes as input the target image I as well as the rendered output of
 73 previous drawing commands. Intuitively, the network looks at the image it wants to explain, as well
 74 as what it has already drawn. It then decides either to stop drawing or proposes another drawing
 75 command to add to the execution trace; if it decides to continue drawing, the predicted primitive is
 76 rendered to its “canvas” and the process repeats.

77 Figure 2 illustrates this architecture. We first pass a 256×256 target image and a rendering of the
 78 trace so far to a convolutional network – these two inputs are represented as separate channels for
 79 the convnet. Given the features extracted by the convnet, a multilayer perceptron then predicts a
 80 distribution over the next drawing command to add to the trace. We predict the drawing command
 81 token-by-token, and condition each token both on the image features and on the previously generated
 82 tokens. For example, the network first decides to emit the `circle` token conditioned on the image
 83 features, then it emits the x coordinate of the circle conditioned on the image features and the `circle`
 84 token, and finally it predicts the y coordinate of the circle conditioned on the image features, the
 85 `circle` token, and the x coordinate. [There are some more details that are important to provide
 86 about this architecture, though possibly in an Appendix: the functional form(s) of the probability
 87 distributions over tokens, the network layer sizes, which MLPs share parameters, etc.]

88 [Planning to move the description of SMC / beam search up here, too?]

89 The distribution over the next drawing command factorizes:

$$\mathbb{P}_\theta[t_1 t_2 \cdots t_K | I, T] = \prod_{k=1}^K \mathbb{P}_\theta[t_k | f_\theta(I, \text{render}(T)), \{t_j\}_{j=1}^{k-1}] \quad (1)$$

90 where $t_1 t_2 \cdots t_K$ are the tokens in the drawing command, I is the target image, T is an execution trace,
 91 θ are the parameters of the neural network, and $f_\theta(\cdot, \cdot)$ is the image feature extractor (convolutional
 92 network). The distribution over execution traces factorizes as:

$$\mathbb{P}_\theta[T | I] = \prod_{n=1}^{|T|} \mathbb{P}_\theta[T_n | I, T_{1:(n-1)}] \times \mathbb{P}_\theta[\text{STOP} | I, T] \quad (2)$$

93 where $|T|$ is the length of execution trace T , and the `STOP` token is emitted by the network to signal
 94 that the execution trace explains the image.

95 We train the network by sampling execution traces T and target images I for randomly generated
 96 scenes, and maximizing (2) wrt θ by gradient ascent. Training does not require backpropagation across
 97 the entire sequence of drawing commands: drawing to the canvas ‘blocks’ the gradients, effectively
 98 offloading memory to an external visual store. In a sense, this model is like an autoregressive variant
 99 of AIR [1] without attention.

100 [I like that you make this connection, but it could be made more precisely. Specifically, (1) the
 101 architecture isn’t *really* recurrent (it uses no hidden state cells), so it’d be good to use a different term
 102 or drop this part of the point: (2) training of recurrent nets is also typically fully-supervised (Most
 103 RNNs lack latent variables per timestep)—if you’re thinking about AIR specifically, maybe just say
 104 that, and (3) it’s like an autoregressive AIR *without attention*.] [Something related to this that’s also
 105 cool to point out: training this model doesn’t require backpropagation across the entire sequence of
 106 drawing commands (drawing to the canvas ‘blocks’ the gradients, effectively offloading memory to
 107 an external (visual) store, so in principle it might be scalable to much longer sequences.]

108 This network suffices to “derender” images like those shown in Figure 3. We can perform a beam
 109 search decoding to recover what the network thinks is the most likely execution trace for images
 110 like these. But, if the network makes a mistake (predicts an incorrect line of code), it has no way
 111 of recovering from the error. In order to derender an image with n objects, it must correctly predict
 112 n drawing commands – so its probability of success will decrease exponentially in n , assuming it
 113 has any nonzero chance of making a mistake. For added robustness as n becomes large, we treat the
 114 neural network outputs as proposals for a SMC sampling scheme. For the SMC sampler, we use pixel
 115 wise distance as a surrogate for a likelihood function; see supplement. Figure 4 compares the neural

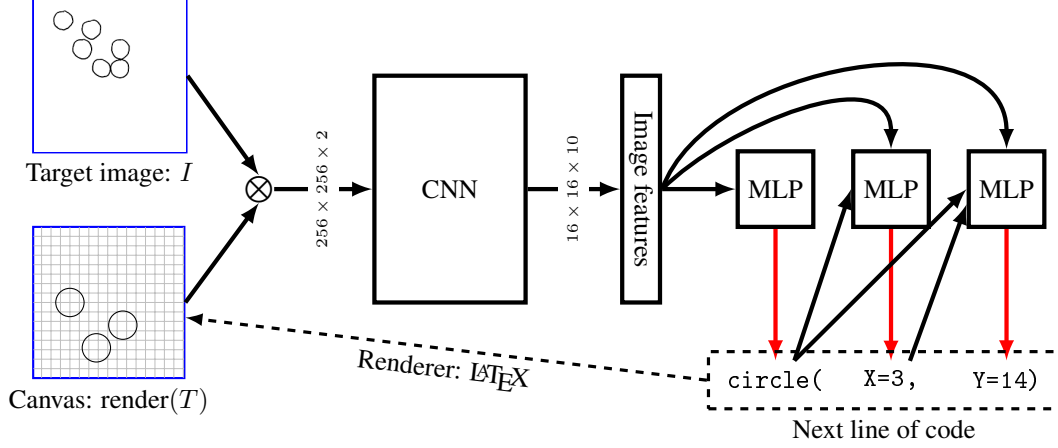


Figure 2: Our neural architecture for inferring the execution trace of a graphics program from its output. **Blue**: network inputs. **Black**: network operations. **Red**: samples from a multinomial. Typewriter font: network outputs. Renderers snapped to a 16×16 grid, illustrated in gray. [Thoughts on improving this figure: (1) Convnet diagrams typically show the sequence of layers, if possible (space might not permit it here, but those thin arrows just aren't doing it for me). (2) Are the target image / canvas convolved down independently, or jointly (i.e. starting as a 2-channel image)? That's an important detail that's not clear with the current figure/explanation. (3) The three circles downstream from 'Image Features' are supposed to be MLPs, I assume(?), but it took me a little while to parse that. Having some visual way of clearly separating network operations from data (color, perhaps) would go a long way.]

<code>circle(x, y)</code>	Circle at (x, y)
<code>rectangle(x_1, y_1, x_2, y_2)</code>	Rectangle with corners at (x_1, y_1) & (x_2, y_2)
<code>LINE(x_1, y_1, x_2, y_2,</code> <code>arrow $\in \{0, 1\}$, dashed $\in \{0, 1\}$)</code>	Line from (x_1, y_1) to (x_2, y_2) , optionally with an arrow and/or dashed
<code>STOP</code>	Finishes execution trace inference

Table 1: The deep network in (2) predicts drawing commands, shown above.

116 network with SMC against the neural network by itself or SMC by itself. Only the combination of the
 117 two passes a critical test of generalization: when trained on images with ≤ 8 objects, it successfully
 118 parses scenes with many more objects than the training data.

119 3.1 Generalizing to hand drawings

120 We believe that converting synthetic, noiseless
 121 images into a restricted subset of $L^A_T_E_X$ has limited
 122 usefulness. A more practical application is
 123 one that extends to hand drawings. We train the
 124 model to generalize to hand drawings by intro-
 125 ducing noise into the renderings of the training
 126 target images. We designed this noise process to
 127 introduce the kinds of variations found in hand
 128 drawings (figure 6). We drew 100 figures by
 129 hand; see figure ???. These were drawn reason-
 130 ably carefully but not perfectly. Because our
 131 model assumes that objects are snapped to a
 132 16×16 grid, we made the drawings on graph paper.




Figure 3: Network is trained to infer execution traces for figures like the three shown above.




Figure 4: Using the model to parse latex output. The model is trained on diagrams with up to 8 objects. As shown above it generalizes to scenes with many more objects. Neither the stochastic search nor the neural network are sufficient on their own.




Figure 5: (a): a hand drawing. (b): Rendering of the parse our model infers for (a). We can generalize to hand drawings like these because we train the model on images corrupted by a noise process designed to resemble the kind of noise introduced by hand drawings - see (c) & (d) for noisy renderings of (b).

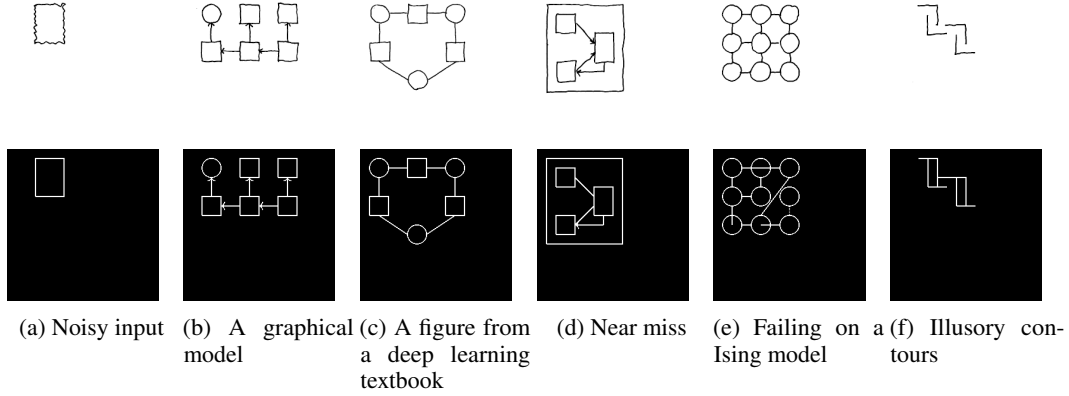


Figure 6: Example drawings above model outputs. See also Fig. 1

133 4 Synthesizing graphics programs from execution traces

134 Although the execution trace of a graphics program describes the parts of a scene, it fails to encode
 135 higher-level features of the image, such as repeated motifs, symmetries or reflections. A *graphics*
 136 *program* better describe structures like these, and we now take as our goal to synthesize simple
 137 graphics programs from their execution traces.

138 We constrain the space of allowed programs by writing down a context free grammar over a space of
 139 programs. Although it might be desirable to synthesize programs in a Turing-complete language like
 140 Lisp or Python, a more tractable approach is to specify what in the program languages community
 141 is called a Domain Specific Language (DSL). Our DSL (Table 2) encodes prior knowledge of what
 142 graphics programs tend to look like.

Program	→	Command; ...; Command
Command	→	circle(Expression, Expression)
Command	→	rectangle(Expression, Expression, Expression, Expression)
Command	→	LINE(Expression, Expression, Expression, Expression, Boolean, Boolean)
Command	→	for($0 \leq \text{Var} < \text{Expression}$) { Program }
Command	→	REFLECT(Axis) { Program }
Expression	→	$Z * \text{Var} + Z$
Var	→	A free (unused) variable
Z	→	an integer
Axis	→	$X = Z$
Axis	→	$Y = Z$

Table 2: Grammar over graphics programs. We allow loops (for), vertical/horizontal reflections (REFLECT), and affine transformations ($Z * \text{Var} + Z$).

143 Given the DSL and a trace T , we want a program that evaluates to T and also minimizes some
 144 measure of program cost:

$$\text{program}(T) = \underset{\substack{p \in \text{DSL} \\ p \text{ evaluates to } T}}{\arg \min} \text{cost}(p) \quad (3)$$


145 An intuitive measure of program cost is its length. We define the cost of a program to be the number
 146 of statements it contains, where a statement is a “Command” in Table 2. [The flow here is a bit
 147 off/backwards. “We want a program that evaluates to T and also minimizes some measure of program
 148 cost”—why do we care about cost? It’d be better to start by making a “Bayesian Occam’s razor”
 149 appeal (e.g. the most compact/general program is the more likely explanation) and then saying that
 150 one way to do this is to minimize a cost function which is proportional to program length.]

151 The constrained optimization problem in equation 3 is intractable in general, but there exist efficient-
 152 in-practice tools for finding exact solutions to program synthesis problems like these. We use the
 153 state-of-the-art Sketch tool [?]. Describing Sketch’s program synthesis algorithm is beyond the
 154 scope of this paper; see supplement. At a high level, Sketch takes as input a space of programs,
 155 along with a specification of the program’s behavior and optionally a cost function. It translates the
 156 synthesis problem into a constraint satisfaction problem, and then uses a quasibootlean solver to find a
 157 minimum cost program satisfying the specification. In exchange for not having any guarantees on
 158 how long it will take to find a minimum cost solution, it comes with the guarantee that it will always
 159 find a globally optimal program.

160 Why synthesize a graphics program, if the execution trace already suffices to recover the objects in an
 161 image? Within our domain of hand-drawn figures, graphics program synthesis has several important
 162 uses:

163 4.1 Extrapolating figures

164 Having access to the source code of a graphics program facilitates coherent, high-level edits to the
 165 figure generated by that program. For example, we can change all of the circles to squares, were
 166 make all of the lines be dashed. We can also **extrapolate** figures by increasing the number of times
 167 that loops are executed.






172




173



174




175



176



177




178



179


180



181




182



183




184



185




186




187







188




189



191



193




195

196




197



198




199



200



201



202




203



204




205



206



207




208



209




210




211


212



213




214




215




216



217




218



219




220




221




222



223




224




225




226



227



228



229




230


231



232




233




234

235






236



237




238



239




240



241




242




243

244



`def plot_figures(figures, axes):
 for figure in figures:
 axes.imshow(figure, cmap=cm.gray)
 axes.set_title(figure.get_filename())
 axes.tight_layout()
 return axes`

245 4.2 Modeling similarity between figures



246

247 4.3 Correcting errors made by the neural network

248 [Seems like you're still fleshing this part out, but I'll give my feedback anyway: (1) This subsection
249 could really use a motivational introduction, e.g. "The program synthesizer can help correct
250 errors/bad proposals from the neural network by favoring execution traces which lead to more
251 concise/general programs." (2) The image likelihood function should probably be introduced sooner,
252 when you talk about SMC/beam search. (3) Where does θ come from? Is it set by hand? Learned?
253 (4) How does Equation 4 get used? Is this a modification to the beam search objective / SMC
254 posterior? If so, it'd be great to have set up the version without it in an earlier section, and then be
255 able to refer to this as a small modification of the previous equation.]


256 Let $L(\cdot|\cdot) : \text{image}^2 \rightarrow \mathcal{R}$ be our likelihood function: it takes two images, an observed target image
257 and a hypothesized program output, and gives the likelihood of the observed image conditioned on
258 the program output. Write $\hat{T}(I)$ for the trace the model predicts for image I .

259 We can extract a few basic features of a program, like its size or how many loops it has, and use these
 260 features to help predict whether a trace is the correct explanation for an image.

$$\hat{T}(I) = \arg \max_T L(I|\text{render}(T)) + \theta \cdot \phi(\text{program}(T)) \quad (4)$$


261 where $\phi(\cdot)$ is a feature extractor for programs. This is equivalent to doing MAP inference in a
 262 generative model where the program is first drawn from a log linear distribution
 263 $\propto \exp(\theta \cdot \phi(\text{program}))$, than the program is executed deterministically, and then we observe a noisy
 264 version of the program's output, where L is the noise model.

267




```
line(6,2,6,3,arrow = False,solid = True)
line(6,2,3,2,arrow = True,solid = True)
rectangle(0,0,8,9)
rectangle(5,3,7,6)
reflect(y = 9)
line(3,7,5,5,arrow = True,solid = True)
rectangle(1,6,3,8)
```

268




```
rectangle(4,2,6,5)
reflect(y = 7)
line(2,6,4,4,arrow = True,solid = True)
rectangle(0,5,2,7)
```

269




```
circle(10,5)
line(3,2,5,4,arrow = True,solid = True)
for (2)
  line(-4*i + 7,3*i + 5,-4*i + 9,5,arrow = True,solid = True)
  rectangle(-5*i + 5,-3*i + 3,5*i + 7,3*i + 7)
  rectangle(1,6*i + 1,3,6*i + 3)
```




270

```
circle(10,4)
line(10,1,10,3,arrow = False,solid = False)
line(10,1,2,1,arrow = True,solid = False)
rectangle(4,2,7,6)
    reflect(y = 8)
    line(2,1,4,3,arrow = True,solid = True)
    line(7,4,9,4,arrow = True,solid = True)
    rectangle(0,0,2,2)
```




271

```
line(12,9,12,0,arrow = True,solid = True)
for (2)
    rectangle(3*i + 6,-2*i + 5,3*i + 8,9)
    rectangle(3*i,7,3*i + 2,9)
```



272

```
line(0,4,3,4,arrow = False,solid = True)
rectangle(2,0,5,3)
for (2)
    line(3*i,3*i + 1,3*i + 2,3*i,arrow = False,solid = True)
    line(0,-3*i + 4,-2*i + 2,3,arrow = False,solid = True)
```




273

```
for (3)
  circle(3*i + 1,6)
  circle(-3*i + 7,1)
  line(-3*i + 7,-1*i + 4,-3*i + 7,5,arrow = False,solid = True)
```



274

```
line(0,0,0,4,arrow = False,solid = True)
```



275

```
line(6,0,0,0,arrow = True,solid = True)
```




276

```
rectangle(0,0,3,4)
```




277

```
circle(1,1)
```




278

```
for (2)
  circle(-5*i + 6,1)
  line(2,5*i + 1,5,5*i + 1,arrow = False,solid = True)
  line(5*i + 1,2,5*i + 1,5,arrow = False,solid = True)
  rectangle(-5*i + 5,5,-5*i + 7,7)
```



279

```
line(2,1,4,1,arrow = True,solid = True)
line(5,0,3,0,arrow = True,solid = True)
line(0,3,2,3,arrow = True,solid = True)
line(3,2,1,2,arrow = True,solid = True)
```




280

```
rectangle(6,0,7,1)
for (3)
  rectangle(-2*i + 4,2*i + 2,-2*i + 5,2*i + 3)
  rectangle(-2*i + 4,2*i,-2*i + 5,2*i + 1)
```




281

```
line(3,0,5,0,arrow = False,solid = True)
line(2,1,4,1,arrow = False,solid = False)
line(0,3,2,3,arrow = False,solid = False)
line(1,2,3,2,arrow = False,solid = True)
```




282

```
circle(9,1)
for (3)
  circle(2*i + 3,-3*i + 10)
  circle(2*i + 1,-3*i + 7)
  line(-2*i + 6,3*i + 1,-2*i + 7,3*i + 3,arrow = False,solid
  line(-2*i + 7,3*i + 3,-2*i + 8,3*i + 1,arrow = False,solid
```



283

```
rectangle(8,0,10,2)
for (3)
  line(-2*i + 7,3*i + 3,-2*i + 8,3*i + 1,arrow = True,solid
  line(2*i + 3,-3*i + 9,2*i + 2,-3*i + 7,arrow = True,solid
  rectangle(-2*i + 6,3*i + 3,-2*i + 8,3*i + 5)
  rectangle(-2*i + 4,3*i,-2*i + 6,3*i + 2)
```


284

```
for (2)
  circle(-1*i + 4,-2*i + 5)
  circle(-6*i + 7,3*i + 3)
  circle(-2*i + 7,1)
```




285

```
line(4,4,2,2,arrow = True,solid = True)
rectangle(3,4,5,6)
rectangle(0,0,2,2)
```




286

```
rectangle(0,4,2,6)
for (2)
  circle(4*i + 5,4*i + 1)
  line(4*i + 4,5,4*i + 2,5,arrow = True,solid = True)
  reflect(x = 10)
  circle(4*i + 5,-4*i + 5)
  line(4*i + 5,4,4*i + 5,2,arrow = True,solid = True)
```




287

```
circle(1,5)
for (2)
  line(-4*i + 8,1,-4*i + 6,1,arrow = True,solid = True)
  rectangle(4*i + 4,4*i,4*i + 6,4*i + 2)
  reflect(x = 10)
  line(-4*i + 5,2,-4*i + 5,4,arrow = True,solid = True)
  rectangle(4*i,4*i,4*i + 2,4*i + 2)
```




288

```
line(4,5,2,5,arrow = True,solid = True)
line(8,5,6,5,arrow = True,solid = True)
for (3)
  line(4*i + 1,4,4*i + 1,2,arrow = True,solid = True)
  rectangle(4*i,0,4*i + 2,2)
  rectangle(-4*i + 8,4,-4*i + 10,6)
```



289

```
for (2)
  line(-4*i + 8,1,-4*i + 6,1,arrow = True,solid = True)
  reflect(x = 10)
  circle(5,-4*i + 5)
  circle(-8*i + 9,4*i + 1)
  line(-4*i + 9,2,-4*i + 9,4,arrow = True,solid = True)
```



290

```
line(5,7,5,6,arrow = True,solid = True)
rectangle(2,0,4,2)
for (2)
    line(2*i + 1,-4*i + 7,2*i + 1,-4*i + 6,arrow = True,solid
    rectangle(0,4*i + 3,6,3*i + 6)
```



291

```
line(3,1,2,1,arrow = True,solid = True)
rectangle(6,0,8,2)
    reflect(x = 5)
    line(6,1,5,1,arrow = True,solid = True)
    rectangle(3,0,5,2)
```



292

```
line(1,3,1,4,arrow = False,solid = True)
line(1,3,1,2,arrow = True,solid = True)
line(1,8,1,6,arrow = True,solid = True)
for (3)
    circle(1,4*i + 1)
```



293

```
line(0,1,1,0,arrow = False,solid = True)
line(1,2,2,1,arrow = False,solid = True)
```




294

```
line(0,2,2,2,arrow = False,solid = True)
line(0,0,0,2,arrow = False,solid = True)
```




295

```
for (3)
  line(0,-1*i + 6,2*i + 2,-1*i + 6,arrow = False,solid = True)
  line(-1*i + 2,2*i,-1*i + 2,4,arrow = False,solid = True)
```




296

```
rectangle(0,0,2,2)
for (2)
  circle(1,-3*i + 4)
  circle(5,-2*i + 4)
  rectangle(-4*i + 4,2*i + 1,-4*i + 6,5)
```




297

```
circle(4,3)
for (3)
    circle(-3*i + 7,5)
    circle(7,-2*i + 5)
    rectangle(-3*i + 6,2*i,-3*i + 8,6)
```



298

```
circle(5,5)
line(2,5,4,5,arrow = False,solid = True)
rectangle(0,4,2,6)
rectangle(0,0,5,3)
```




299

```
line(0,3,6,3,arrow = False,solid = False)
line(0,0,0,3,arrow = False,solid = True)
line(0,0,6,0,arrow = False,solid = False)
line(6,0,6,3,arrow = False,solid = True)
```




300

```
for (2)
  line(-1*i + 2,-1*i + 3,-1*i + 2,-1*i + 5,arrow = False,solid = True)
  line(3,-3*i + 3,5,-3*i + 3,arrow = False,solid = True)
  line(4*i,-5*i + 5,2*i + 2,-3*i + 5,arrow = False,solid = True)
  line(-4*i + 5,1,-2*i + 5,-1*i + 3,arrow = False,solid = True)
```



301


```
circle(1,6)
for (2)
  circle(-5*i + 6,1)
  line(-1*i + 2,-1*i + 6,-5*i + 6,2,arrow = True,solid = True)
```



302

```
for (2)
  rectangle(-8*i + 8,5*i,-7*i + 9,6*i + 1)
  rectangle(0,-8*i + 8,3*i + 1,-5*i + 9)
  rectangle(-2*i + 7,3*i + 2,9,5*i + 4)
```

303




```
circle(1,8)
for (2)
  line(-4*i + 5,5*i + 2,-2*i + 5,5,arrow = False,solid = True)
  rectangle(4*i,-5*i + 5,6,-7*i + 9)
```



304

```
circle(4,1)
reflect(x = 8)
circle(1,1)
```



305

```
rectangle(0,0,1,6)
rectangle(4,0,5,6)
rectangle(2,0,3,6)
```




306

```
line(4,1,4,5,arrow = False,solid = False)
line(4,0,4,1,arrow = False,solid = False)
line(0,0,0,5,arrow = False,solid = False)
```




307

```
line(0,0,0,5,arrow = False,solid = True)
line(4,0,4,5,arrow = False,solid = True)
```



308


```
reflect(x = 12)
circle(4,1)
line(9,1,10,1,arrow = False,solid = True)
rectangle(0,0,2,2)
```



309


```
circle(7,6)
for (2)
  reflect(y = 12)
  line(2*i + 2,4*i + 2,4*i + 2,2*i + 4,arrow = True,solid = True)
  rectangle(-1*i + 1,-6*i + 10,3,-4*i + 12)
```


310




```
reflect(x = 9)
line(8,3,8,6,arrow = False,solid = True)
  reflect(y = 9)
  circle(1,1)
  line(3,1,6,1,arrow = False,solid = True)
```

311




```
reflect(y = 11)
rectangle(4,9,7,10)
  reflect(x = 11)
  rectangle(8,0,11,3)
  rectangle(9,4,10,7)
```

312




```
for (3)
  line(3,-1*i + 2,5,-1*i + 2,arrow = False,solid = True)
  line(-1*i + 2,3,-1*i + 4,3,arrow = False,solid = True)
```

313



```
rectangle(0,0,3,3)
for (2)
  rectangle(0,-3*i + 7,-1*i + 3,-4*i + 10)
  rectangle(-3*i + 7,0,-3*i + 9,2)
```

314




```
line(8,3,9,3,arrow = False,solid = True)
for (2)
  line(2,-5*i + 6,4,-1*i + 4,arrow = True,solid = True)
  line(7,-1*i + 4,-2*i + 9,-1*i + 4,arrow = True,solid = True)
  rectangle(0,-5*i + 5,2,-5*i + 7)
  rectangle(-5*i + 9,2,-5*i + 12,5)
```

315




```
rectangle(4,4,7,5)
rectangle(0,0,3,1)
rectangle(2,2,5,3)
```

316



```
circle(4,10)
for (3)
  circle(-3*i + 7,5)
  circle(3*i + 1,1)
  line(3*i + 1,4,3*i + 1,2,arrow = True,solid = True)
  line(4,9,-3*i + 7,6,arrow = True,solid = True)
```

317




```
line(0,8,8,8,arrow = False,solid = True)
line(6,8,6,4,arrow = True,solid = True)
line(2,8,2,6,arrow = True,solid = True)
line(4,8,4,0,arrow = True,solid = True)
```

318




```
line(2,3,2,5,arrow = False,solid = True)
rectangle(1,5,3,7)
rectangle(0,0,4,7)
rectangle(1,1,3,3)
```

319




```
circle(1,5)
line(1,4,1,2,arrow = True,solid = True)
rectangle(0,0,2,2)
```

320




```
circle(3,1)
for (2)
    reflect(x = 6)
    circle(4*i + 1,-4*i + 8)
    circle(-4*i + 5,5*i + 1)
    rectangle(-4*i + 4,-3*i + 3,6,-7*i + 9)
```

321




```
for (3)
    circle(-4*i + 9,4)
    circle(4*i + 1,7)
    circle(-4*i + 9,1)
```




322

```
line(8,0,8,7,arrow = True,solid = True)
line(8,0,0,0,arrow = True,solid = True)
for (3)
    rectangle(-2*i + 6,0,-2*i + 7,-1*i + 5)
```



323

```
line(4,0,0,0,arrow = False,solid = False)
```




324

```
rectangle(0,3,2,5)
for (2)
    circle(3*i + 4,1)
    circle(-2*i + 7,-3*i + 7)
    line(-1*i + 6,-4*i + 7,2*i + 5,-3*i + 5,arrow = True,solid = True)
    line(3*i + 2,-1*i + 4,4,-2*i + 4,arrow = True,solid = True)
```




325

```
circle(2,1)
circle(6,1)
line(5,1,3,1,arrow = True,solid = True)
rectangle(0,0,7,2)
```




326

```
rectangle(2,1,4,3)
rectangle(5,0,8,3)
rectangle(0,2,1,3)
```




327

```
rectangle(2,2,3,3)
rectangle(0,0,5,5)
rectangle(1,1,4,4)
```




328

```
reflect(x = 6)
line(5,2,5,4,arrow = False,solid = True)
    reflect(y = 6)
        line(2,1,4,1,arrow = False,solid = True)
            rectangle(0,0,2,2)
```




329

```
reflect(x = 6)
line(5,2,5,4,arrow = False,solid = True)
    reflect(y = 6)
        circle(1,1)
            line(2,1,4,1,arrow = False,solid = True)
```



330

```
line(0,2,7,2,arrow = False,solid = True)
line(1,1,6,1,arrow = False,solid = True)
line(2,0,5,0,arrow = False,solid = True)
```




331

```
line(1,5,5,1,arrow = False,solid = True)
line(1,4,5,0,arrow = False,solid = True)
rectangle(0,4,1,5)
rectangle(5,0,6,1)
```




332

```
for (3)
  circle(4*i + 1,1)
  rectangle(-4*i + 8,0,-4*i + 10,2)
```



333

```
reflect(x = 5)
circle(1,1)
line(4,4,4,2,arrow = True,solid = True)
rectangle(0,4,5,6)
```





334

```

circle(3,1)
for (2)
  line(-4*i + 5,8,-4*i + 5,6,arrow = True,solid = True)
  reflect(x = 6)
  circle(4*i + 1,-4*i + 9)
  line(5,-4*i + 8,2,-3*i + 5,arrow = True,solid = True)

```




335

```

reflect(x = 8)
circle(4,4)
  reflect(y = 8)
  circle(7,7)
  rectangle(2,2,3,3)

```




336

```


rectangle(0,0,6,6)
  reflect(x = 6)
  rectangle(2,2,4,4)
  for (3)
    circle(3,2*i + 1)
    circle(5,-2*i + 5)

```




337

```
for (3)
  line(4*i + 2,5,4*i + 4,5,arrow = True,solid = True)
  reflect(x = 10)
  circle(-4*i + 13,1)
  circle(4*i + 5,5)
  line(4*i + 5,4,4*i + 5,2,arrow = True,solid = True)
```




338

```
line(4,5,4,10,arrow = False,solid = True)
  reflect(x = 8)
  circle(4,1)
  circle(1,8)
  line(4,5,8,2,arrow = False,solid = True)
```



339

```
reflect(x = 10)
line(6,8,8,8,arrow = False,solid = True)
  for (2)
    circle(-4*i + 5,7*i + 1)
    line(3*i + 6,4*i + 1,9,4*i + 3,arrow = False,solid = T
    rectangle(4*i + 4,-4*i + 7,4*i + 6,-4*i + 9)
```




340

```
line(0,6,12,6,arrow = False,solid = True)
line(6,6,6,5,arrow = True,solid = True)
line(8,3,7,4,arrow = True,solid = True)
line(3,2,5,4,arrow = True,solid = True)
    line(0,8,8,0,arrow = False,solid = True)
```



341

```
rectangle(9,0,11,2)
rectangle(6,0,8,2)
    reflect(x = 5)
    circle(4,1)
```



342


```
for (2)
    line(-4*i + 7,6,-6*i + 8,5,arrow = True,solid = True)
        reflect(x = 10)
        circle(2*i + 7,7)
        circle(3*i + 2,-3*i + 4)
        line(7*i + 1,-3*i + 6,3*i + 2,-3*i + 5,arrow = True,so
```

343




```
circle(1,1)
for (3)
  circle(1,-3*i + 10)
  line(1,-3*i + 9,1,-3*i + 8,arrow = True,solid = True)
```

344



```
for (2)
  rectangle(-7*i + 7,-7*i + 7,-7*i + 8,-7*i + 8)
  rectangle(-6*i + 6,-2*i + 4,-6*i + 8,-2*i + 6)
  rectangle(-5*i + 5,5*i,-5*i + 8,5*i + 3)
```

345



```
reflect(x = 8)
rectangle(0,0,1,1)
rectangle(0,2,2,4)
rectangle(0,5,3,8)
```

346 References

- 347 [1] SM Eslami, N Heess, and T Weber. Attend, infer, repeat: Fast scene understanding with generative models.
 348 arxiv preprint arxiv:..., 2016. URL <http://arxiv.org/abs/1603.08575>.
- 349 [2] Daniel Ritchie, Anna Thomas, Pat Hanrahan, and Noah Goodman. Neurally-guided procedural models:
 350 Amortized inference for procedural graphics programs using neural networks. In *Advances In Neural*
 351 *Information Processing Systems*, pages 622–630, 2016.
- 352 [3] Brian Hempel and Ravi Chugh. Semi-automated svg programming via direct manipulation. In *Proceedings*
 353 *of the 29th Annual Symposium on User Interface Software and Technology*, UIST '16, pages 379–390, New
 354 York, NY, USA, 2016. ACM.

- 355 [4] Haibin Huang, Evangelos Kalogerakis, Ersin Yumer, and Radomir Mech. Shape synthesis from sketches via
356 procedural models and convolutional networks. *IEEE transactions on visualization and computer graphics*,
357 2017.
- 358 [5] Gen Nishida, Ignacio Garcia-Dorado, Daniel G. Aliaga, Bedrich Benes, and Adrien Bousseau. Interactive
359 sketching of urban procedural models. *ACM Trans. Graph.*, 35(4), 2016.