
Inferring Graphics Programs from Images

Anonymous Author(s)

Affiliation

Address

email

Abstract

1

2 1 Introducing visual programs

3 In this work we consider programs that draw diagrams, similar to those found in papers.

4 We develop a hybrid architecture for inferring graphics programs. Our approach uses a deep network
5 infer an execution trace from an image; this recovers primitive drawing operations like lines, circles,
6 or arrows. For added robustness we use the deep network as a proposal distribution for a stochastic
7 search over execution traces. Finally we use techniques in the program synthesis community to
8 recover the program from its trace.

9 Each of these three components – the deep network, the stochastic search, the program synthesizer
10 – confers its own advantages. From the deep network we get a very fast system that can recover
11 plausible execution traces in about a minute. From the stochastic search we get added robustness;
12 essentially the stochastic search can correct mistakes made by the deep network’s proposals. From
13 the program synthesizer we get abstraction: our system recovers coordinate transformations, for
14 loops, and subroutines, which are useful for downstream tasks.

15 2 Neural architecture for inferring image parses

16 We developed a deep network architecture for efficiently inferring a execution trace, T , from an
17 image, I . Our model constructs the trace one drawing command at a time. When predicting the next
18 drawing command, the network takes as input the target image I as well as the rendered output of
19 previous drawing commands. Intuitively, the network looks at the image it wants to explain, as well
20 as what it has already drawn. It then decides either to stop drawing or proposes another drawing
21 command to add to the execution trace; if it decides to continue drawing, the predicted primitive is
22 rendered to its “canvas” and the process repeats.

23 Figure 2 illustrates this architecture. We first pass the target image and a rendering of the trace so far
24 to a convolutional network. Given the features extracted by the convolutional network, a multilayer
25 perceptron then predicts a distribution over the next drawing command to add to the trace. We predict
26 the drawing command token-by-token, and condition each token both on the image features and on
27 the previously generated tokens. For example, the network first decides to emit the CIRCLE token
28 conditioned on the image features, then it emits the x coordinate of the circle conditioned on the
29 image features and the CIRCLE token, and finally it predicts the y coordinate of the circle conditioned
30 on the image features, the CIRCLE token, and the x coordinate.

31 The distribution over the next drawing command factorizes:

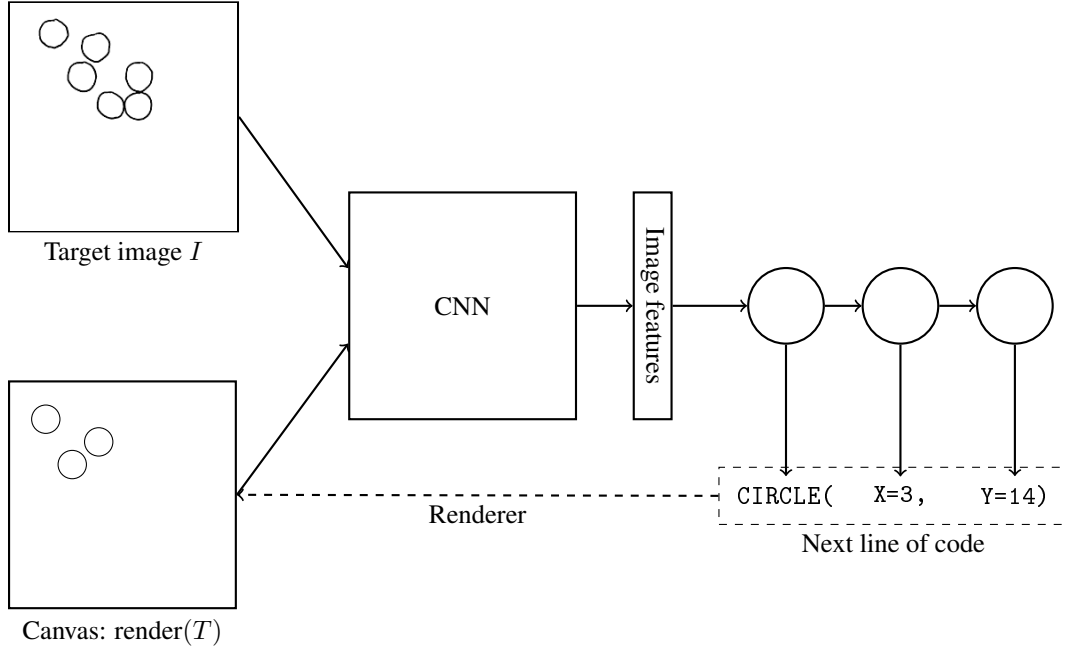
$$\mathbb{P}_\theta[t_1 t_2 \cdots t_K | I, T] = \prod_{k=1}^K \mathbb{P}_\theta[t_k | f_\theta(I, \text{render}(T)), \{t_j\}_{j=1}^{k-1}] \quad (1)$$

where $t_1 t_2 \dots t_K$ are the tokens in the drawing command, I is the target image, T is an execution trace, θ are the parameters of the neural network, and $f_\theta(\cdot, \cdot)$ is the image feature extractor (convolutional network). The distribution over execution traces factorizes as:

$$\mathbb{P}_\theta[T|I] = \prod_{n=1}^{|T|} \mathbb{P}_\theta[T_n|I, T_{1:(n-1)}] \times \mathbb{P}_\theta[\text{STOP}|I, T] \quad (2)$$

where $|T|$ is the length of execution trace T , and the STOP token is emitted by the network to signal that the execution trace explains the image.

We train the network by sampling execution traces T and target images I for randomly generated scenes, and maximizing (2) wrt θ by gradient ascent. Despite the architecture being recurrent, training is fully supervised. In a sense, this model is like an autoregressive variant of AIR. When we have the generative model (the rendering function), we need not solve an unsupervised or reinforcement learning problem.



3 Generalizing to hand drawings

4 Neural networks for guiding SMC

Let $L(\cdot) : \text{image}^2 \rightarrow \mathcal{R}$ be our likelihood function: it takes two images, an observed target image and a hypothesized program output, and gives the likelihood of the observed image conditioned on the program output. We want to sample from:

$$\mathbb{P}[p|x] \propto L(x|\text{render}(p))\mathbb{P}[p] \quad (3)$$

where $\mathbb{P}[p]$ is the prior probability of program p , and x is the observed image.

Let p be a program with L lines, which we will write as $p = (p_1, p_2, \dots, p_L)$. Assume the prior factors into:

$$\mathbb{P}[p] \propto \prod_{l \leq L} \mathbb{P}[p_l] \quad (4)$$

Define the distribution $q_L(\cdot)$, which happens to be proportional to the above posterior:

$$q_L(p_1, p_2, \dots, p_{L-1}, p_L) \propto q_{L-1}(p_1, p_2, \dots, p_{L-1}) \times \frac{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}, p_L))}{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}))} \times \mathbb{P}[p_L] \quad (5)$$




Figure 1: (a): a hand drawing. (b): Rendering of the parse our model infers for (a). We can generalize to hand drawings like these because we train the model on images corrupted by a noise process designed to resemble the kind of noise introduced by hand drawings - see (c) & (d) for noisy renderings of (b).

Now suppose we have some samples from $q_{L-1}(\cdot)$, and that we then sample a p_L from a distribution proportional to $\frac{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}, p_L))}{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}))} \times \mathbb{P}[p_L]$. The resulting programs p are distributed according to q_L , and so are also distributed according to $\mathbb{P}[p|x]$.

How do we sample p_L from a distribution proportional to $\frac{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}, p_L))}{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}))} \times \mathbb{P}[p_L]$? We have a neural network that takes as input the target image x and the program so far, and produces a distribution over next lines of code (p_L). We write $\text{NN}(p_L|p_1, \dots, p_{L-1}; x)$ for the distribution output by the neural network. So we can sample from NN and then weight the samples by:

$$w(p_L) = \frac{\mathbb{P}[p_L]}{\text{NN}(p_L|p_1, \dots, p_{L-1}; x)} \times \frac{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}, p_L))}{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}))} \quad (6)$$

Then we can resample from these now weighted samples to get a new population of particles (here programs are particles), where each program now has L lines instead of $L - 1$.

This procedure can be seen as a particle filter, where each successive latent variable is another line of code, and the emission probabilities are successive ratios of likelihoods under $L(\cdot|\cdot)$.

Comments for Dan. Right now I'm not actually sampling from the neural network - instead, I enumerate the top few hundred lines of code suggested by the network, and then weight them by their likelihoods. So actually the form of NN is:

$$\text{NN}(p_L|p_1, \dots, p_{L-1}; x) \propto \begin{cases} 1, & \text{if } p_L \in \text{top hundred neural network proposals} \\ 0, & \text{otherwise.} \end{cases} \quad (7)$$

Do you think this is a problem? The neural network puts almost all of its mass on a few guesses. In order to get the correct line of code I sometimes need to get something like the 50th top guess, so I don't want to literally just sample from the distribution suggested by the neural network.

5 Submission of papers to NIPS 2017

NIPS requires electronic submissions. The electronic submission site is

<https://cmt.research.microsoft.com/NIPS2017/>

Please read carefully the instructions below and follow them faithfully.

Algorithm 1 Neurally guided SMC

Input: Neural network NN, beam size N , maximum length L , target image x

Output: Samples of the program trace

Set $B_0 = \{\text{empty program}\}$

for $1 \leq l \leq L$ **do**

for $1 \leq n \leq N$ **do**

$p_n \sim \text{Uniform}(B_{l-1})$

$p'_n \sim \text{NN}(\text{render}(p), x)$

 Define $r_n = p'_n \cdot p_n$

 Set $\tilde{w}(r_n) = \frac{L(x|r_n)}{L(x|p_n)} \times \frac{\mathbb{P}[p'_n]}{\mathbb{P}[p'_n = \text{NN}(\text{render}(p), x)]}$

end for

 Define $w(p) = \frac{\tilde{w}(p)}{\sum_{p'} \tilde{w}(p')}$

 Set B_l to be N samples from r_n distributed according to $w(\cdot)$

end for

return $\{p : p \in B_{l \leq L}, p \text{ is finished}\}$

73 5.1 Style

74 Papers to be submitted to NIPS 2017 must be prepared according to the instructions presented here.
75 Papers may only be up to eight pages long, including figures. This does not include acknowledgments
76 and cited references which are allowed on subsequent pages. Papers that exceed these limits will not
77 be reviewed, or in any other way considered for presentation at the conference.

78 The margins in 2017 are the same as since 2007, which allow for $\sim 15\%$ more words in the paper
79 compared to earlier years.

80 Authors are required to use the NIPS L^AT_EX style files obtainable at the NIPS website as indicated
81 below. Please make sure you use the current files and not previous versions. Tweaking the style files
82 may be grounds for rejection.

83 5.2 Retrieval of style files

84 The style files for NIPS and other conference information are available on the World Wide Web at

85 <http://www.nips.cc/>

86 The file `nips_2017.pdf` contains these instructions and illustrates the various formatting require-
87 ments your NIPS paper must satisfy.

88 The only supported style file for NIPS 2017 is `nips_2017.sty`, rewritten for L^AT_EX 2 ϵ . **Previous**
89 **style files for L^AT_EX 2.09, Microsoft Word, and RTF are no longer supported!**

90 The new L^AT_EX style file contains two optional arguments: `final`, which creates a camera-ready copy,
91 and `nonatbib`, which will not load the `natbib` package for you in case of package clash.

92 At submission time, please omit the `final` option. This will anonymize your submission and add
93 line numbers to aid review. Please do *not* refer to these line numbers in your paper as they will be
94 removed during generation of camera-ready copies.

95 The file `nips_2017.tex` may be used as a “shell” for writing your paper. All you have to do is
96 replace the author, title, abstract, and text of the paper with your own.

97 The formatting instructions contained in these style files are summarized in Sections 6, 7, and 8
98 below.

99 6 General formatting instructions

100 The text must be confined within a rectangle 5.5 inches (33 picas) wide and 9 inches (54 picas) long.
101 The left margin is 1.5 inch (9 picas). Use 10 point type with a vertical spacing (leading) of 11 points.

102 Times New Roman is the preferred typeface throughout, and will be selected for you by default.
103 Paragraphs are separated by $\frac{1}{2}$ line space (5.5 points), with no indentation.

104 The paper title should be 17 point, initial caps/lower case, bold, centered between two horizontal
105 rules. The top rule should be 4 points thick and the bottom rule should be 1 point thick. Allow $\frac{1}{4}$ inch
106 space above and below the title to rules. All pages should start at 1 inch (6 picas) from the top of the
107 page.

108 For the final version, authors' names are set in boldface, and each name is centered above the
109 corresponding address. The lead author's name is to be listed first (left-most), and the co-authors'
110 names (if different address) are set to follow. If there is only one co-author, list both author and
111 co-author side by side.

112 Please pay special attention to the instructions in Section 8 regarding figures, tables, acknowledgments,
113 and references.

114 **7 Headings: first level**

115 All headings should be lower case (except for first word and proper nouns), flush left, and bold.

116 First-level headings should be in 12-point type.

117 **7.1 Headings: second level**

118 Second-level headings should be in 10-point type.

119 **7.1.1 Headings: third level**

120 Third-level headings should be in 10-point type.

121 **Paragraphs** There is also a `\paragraph` command available, which sets the heading in bold, flush
122 left, and inline with the text, with the heading followed by 1 em of space.

123 **8 Citations, figures, tables, references**

124 These instructions apply to everyone.

125 **8.1 Citations within the text**

126 The `natbib` package will be loaded for you by default. Citations may be author/year or numeric, as
127 long as you maintain internal consistency. As to the format of the references themselves, any style is
128 acceptable as long as it is used consistently.

129 The documentation for `natbib` may be found at

130 `http://mirrors.ctan.org/macros/latex/contrib/natbib/natnotes.pdf`

131 Of note is the command `\citet`, which produces citations appropriate for use in inline text. For
132 example,

133 `\citet{hasselmo}` investigated\dots

134 produces

135 Hasselmo, et al. (1995) investigated...

136 If you wish to load the `natbib` package with options, you may add the following before loading the
137 `nips_2017` package:

138 `\PassOptionsToPackage{options}{natbib}`

139 If `natbib` clashes with another package you load, you can add the optional argument `nonatbib`
140 when loading the style file:

141 `\usepackage[nonatbib]{nips_2017}`

142 As submission is double blind, refer to your own published work in the third person. That is, use “In
143 the previous work of Jones et al. [4],” not “In our previous work [4].” If you cite your other papers
144 that are not widely available (e.g., a journal paper under review), use anonymous author names in the
145 citation, e.g., an author of the form “A. Anonymous.”

146 8.2 Footnotes

147 Footnotes should be used sparingly. If you do require a footnote, indicate footnotes with a number¹
148 in the text. Place the footnotes at the bottom of the page on which they appear. Precede the footnote
149 with a horizontal rule of 2 inches (12 picas).

150 Note that footnotes are properly typeset *after* punctuation marks.²

151 8.3 Figures

152 All artwork must be neat, clean, and legible. Lines should be dark enough for purposes of reproduction.
153 The figure number and caption always appear after the figure. Place one line space before the figure
154 caption and one line space after the figure. The figure caption should be lower case (except for first
155 word and proper nouns); figures are numbered consecutively.

156 You may use color figures. However, it is best for the figure captions and the paper body to be legible
if the paper is printed in either black/white or in color.




Figure 2: Sample figure caption.

157

158 8.4 Tables

159 All tables must be centered, neat, clean and legible. The table number and title always appear before
160 the table. See Table 1.

161 Place one line space before the table title, one line space after the table title, and one line space after
162 the table. The table title must be lower case (except for first word and proper nouns); tables are
163 numbered consecutively.

164 Note that publication-quality tables *do not contain vertical rules*. We strongly suggest the use of the
165 `booktabs` package, which allows for typesetting high-quality, professional tables:

166 <https://www.ctan.org/pkg/booktabs>

167 This package was used to typeset Table 1.

¹Sample of the first footnote.

²As in this example.

Table 1: Sample table title

Part		
Name	Description	Size (μm)
Dendrite	Input terminal	~ 100
Axon	Output terminal	~ 10
Soma	Cell body	up to 10^6

9 Final instructions

Do not change any aspects of the formatting parameters in the style files. In particular, do not modify the width or length of the rectangle the text should fit into, and do not change font sizes (except perhaps in the **References** section; see below). Please note that pages should be numbered.

10 Preparing PDF files

Please prepare submission files with paper size “US Letter,” and not, for example, “A4.”

Fonts were the main cause of problems in the past years. Your PDF file must only contain Type 1 or Embedded TrueType fonts. Here are a few instructions to achieve this.

- You should directly generate PDF files using `pdflatex`.
- You can check which fonts a PDF files uses. In Acrobat Reader, select the menu Files>Document Properties>Fonts and select Show All Fonts. You can also use the program `pdffonts` which comes with `xpdf` and is available out-of-the-box on most Linux machines.
- The IEEE has recommendations for generating PDF files whose fonts are also acceptable for NIPS. Please see <http://www.emfield.org/icuwb2010/downloads/IEEE-PDF-SpecV32.pdf>
- `xfig` “patterned” shapes are implemented with bitmap fonts. Use “solid” shapes instead.
- The `\bbold` package almost always uses bitmap fonts. You should use the equivalent AMS Fonts:

```
\usepackage{amsfonts}
```

followed by, e.g., `\mathbb{R}`, `\mathbb{N}`, or `\mathbb{C}` for \mathbb{R} , \mathbb{N} or \mathbb{C} . You can also use the following workaround for reals, natural and complex:

```
\newcommand{\RR}{I\!\!R} %real numbers
\newcommand{\Nat}{I\!\!N} %natural numbers
\newcommand{\CC}{I\!\!C} %complex numbers
```

Note that `amsfonts` is automatically loaded by the `amssymb` package.

If your file contains type 3 fonts or non embedded TrueType fonts, we will ask you to fix it.

10.1 Margins in L^AT_EX

Most of the margin problems come from figures positioned by hand using `\special` or other commands. We suggest using the command `\includegraphics` from the `graphicx` package. Always specify the figure width as a multiple of the line width as in the example below:

```
\usepackage[pdftex]{graphicx} ...
\includegraphics[width=0.8\linewidth]{myfile.pdf}
```

See Section 4.4 in the `graphics` bundle documentation (<http://mirrors.ctan.org/macros/latex/required/graphics/grfguide.pdf>)

A number of width problems arise when L^AT_EX cannot properly hyphenate a line. Please give LaTeX hyphenation hints using the `\-` command when necessary.

204 **Acknowledgments**

205 Use unnumbered third level headings for the acknowledgments. All acknowledgments go at the end
206 of the paper. Do not include acknowledgments in the anonymized submission, only in the final paper.

207 **References**

208 References follow the acknowledgments. Use unnumbered first-level heading for the references. Any
209 choice of citation style is acceptable as long as you are consistent. It is permissible to reduce the font
210 size to small (9 point) when listing the references. **Remember that you can go over 8 pages as**
211 **long as the subsequent ones contain *only* cited references.**

212 [1] Alexander, J.A. & Mozer, M.C. (1995) Template-based algorithms for connectionist rule extraction. In
213 G. Tesauro, D.S. Touretzky and T.K. Leen (eds.), *Advances in Neural Information Processing Systems 7*, pp.
214 609–616. Cambridge, MA: MIT Press.

215 [2] Bower, J.M. & Beeman, D. (1995) *The Book of GENESIS: Exploring Realistic Neural Models with the*
216 *GENeral NEural Simulation System*. New York: TELOS/Springer-Verlag.

217 [3] Hasselmo, M.E., Schnell, E. & Barkai, E. (1995) Dynamics of learning and recall at excitatory recurrent
218 synapses and cholinergic modulation in rat hippocampal region CA3. *Journal of Neuroscience* **15**(7):5249-5262.