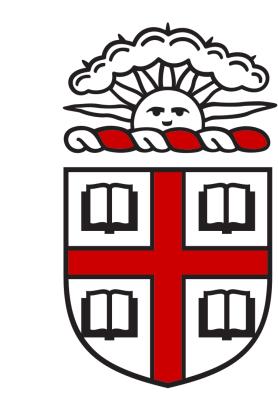
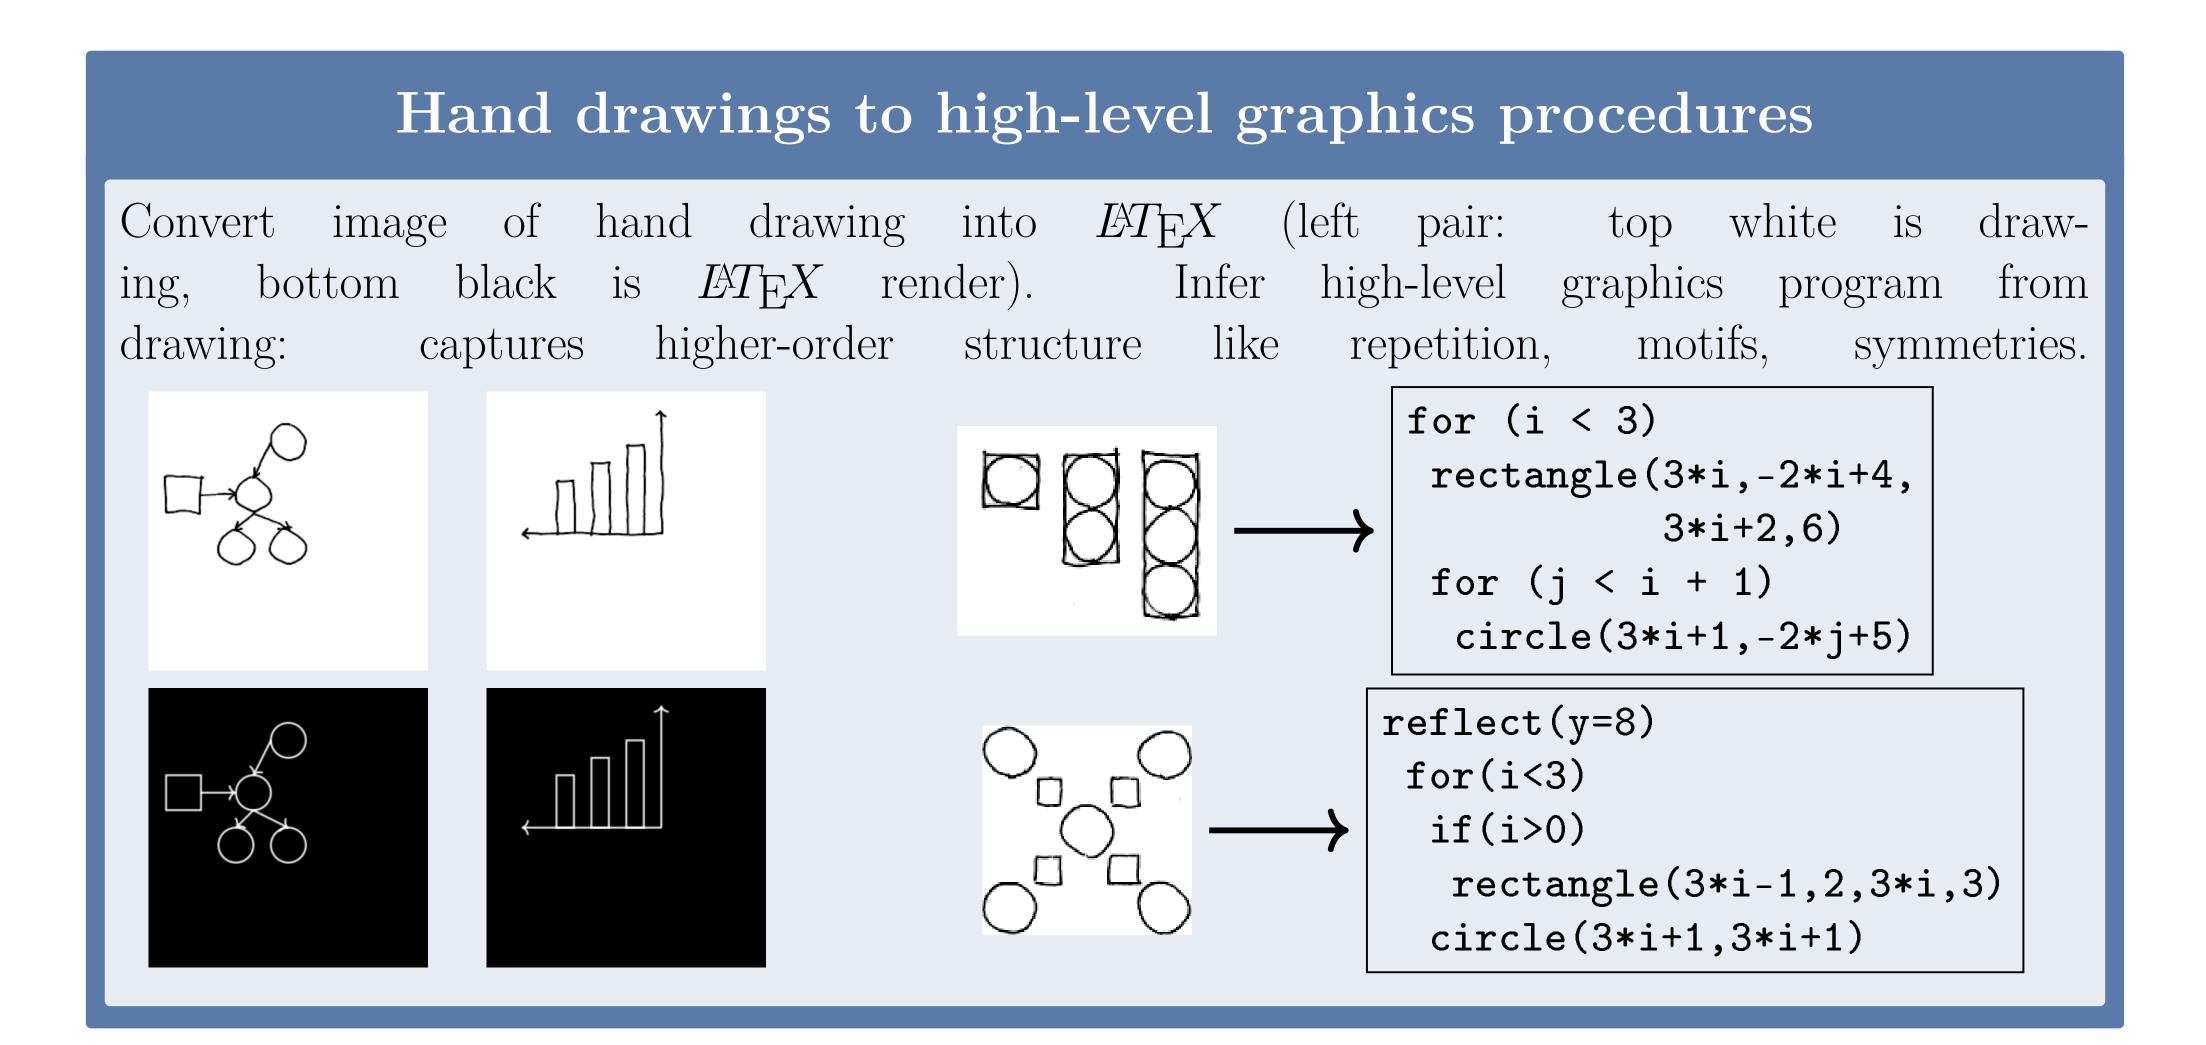


Learning to Infer Graphics Programs from Hand-Drawn Images

Kevin Ellis, Daniel Ritchie, Armando Solar-Lezama, Joshua B. Tenenbaum Massachusetts Institute of Technology & Brown University





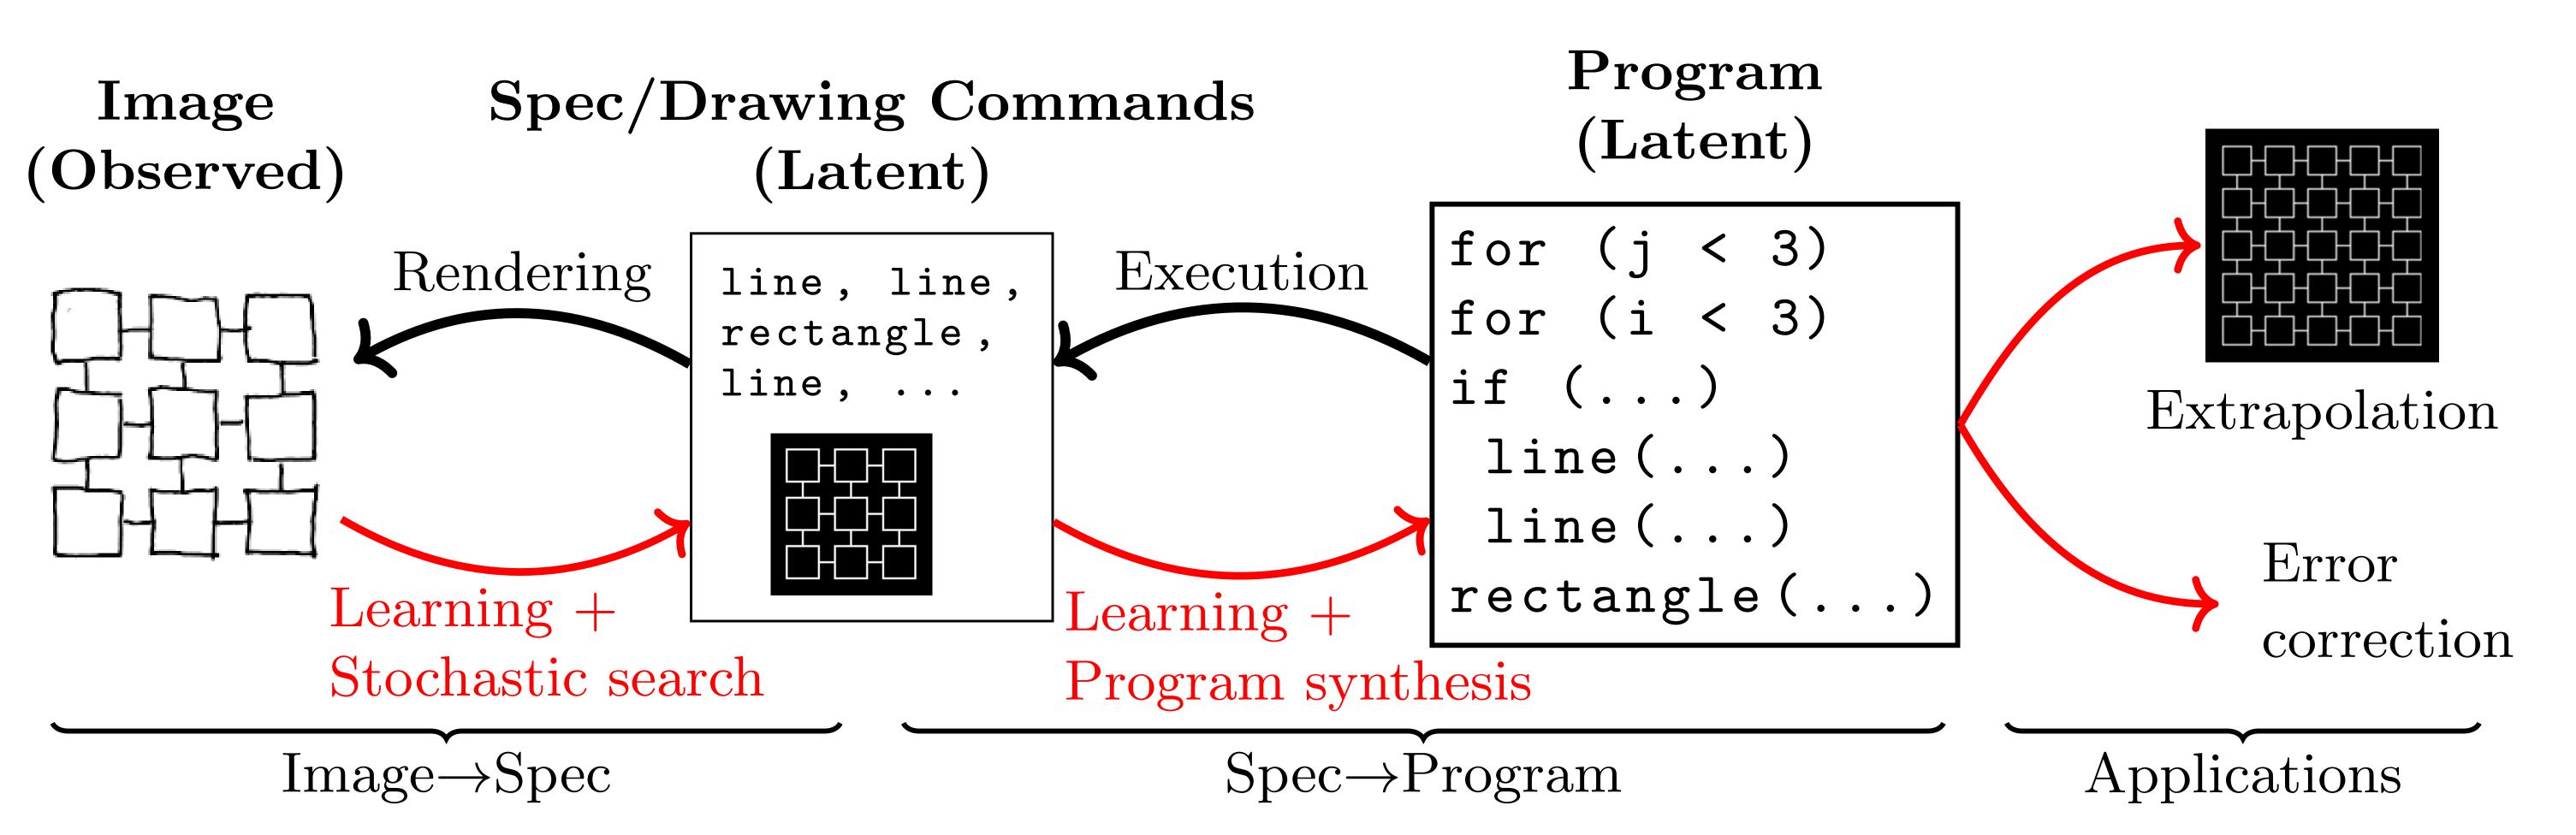
Example System Outputs

Drawing	Spec	Program	Compression factor
	Line(2,15, 4,15) Line(4,9, 4,13) Line(3,11, 3,14) Line(2,13, 2,15) Line(3,14, 6,14) Line(4,13, 8,13)	<pre>for(i<3) line(i,-1*i+6,</pre>	$\frac{6}{3} = 2x$
	Line(5,13,2,10,arrow) Circle(5,9) Circle(8,5) Line(2,8, 2,6,arrow) Circle(2,5) etc; 13 lines	<pre>circle(4,10) for(i<3) circle(-3*i+7,5) circle(-3*i+7,1) line(-3*i+7,4,-3*i+7,2,arrolline(4,9,-3*i+7,6,arrow)</pre>	$\frac{13}{6} = 2.2 \mathrm{x}$
9-0-0-0-0-0-0-0-0-0-0-0-0-0-0-0-0-0-0-0	Circle(5,8) Circle(2,8) Circle(8,11) Line(2,9, 2,10) Circle(8,8) Line(3,8, 4,8) Line(3,11, 4,11) etc; 21 lines	<pre>for(i<3) for(j<3) if(j>0) line(-3*j+8,-3*i+7,</pre>	$\frac{21}{6} = 3.5x$
	Rectangle (1,10,3,11) Rectangle (1,12,3,13) Rectangle (4,8,6,9) Rectangle (4,10,6,11) etc; 16 lines	for(i<4) for(j<4) rectangle(-3*i+9,-2*j+6, -3*i+11,-2*j+7)	$\frac{16}{3} = 5.3x$
	Line(3,10,3,14,arrow) Rectangle(11,8,15,10) Rectangle(11,14,15,15) Line(13,10,13,14,arrow) etc; 16 lines	<pre>for(i<3) line(7,1,5*i+2,3,arrow) for(j<i+1) if(j="">0) line(5*j-1,9,5*i,5,arrow) line(5*j+2,5,5*j+2,9,arrow) rectangle(5*i,3,5*i+4,5) rectangle(5*i,9,5*i+4,10) rectangle(2,0,12,1)</i+1)></pre>	
	Circle(2,8) Rectangle(6,9, 7,10) Circle(8,8) Rectangle(6,12, 7,13) Rectangle(3,9, 4,10) etc; 9 lines	<pre>reflect(y=8) for(i<3) if(i>0) rectangle(3*i-1,2,3*i,3) circle(3*i+1,3*i+1)</pre>	$\frac{9}{5} = 1.8x$

Acknowledgments

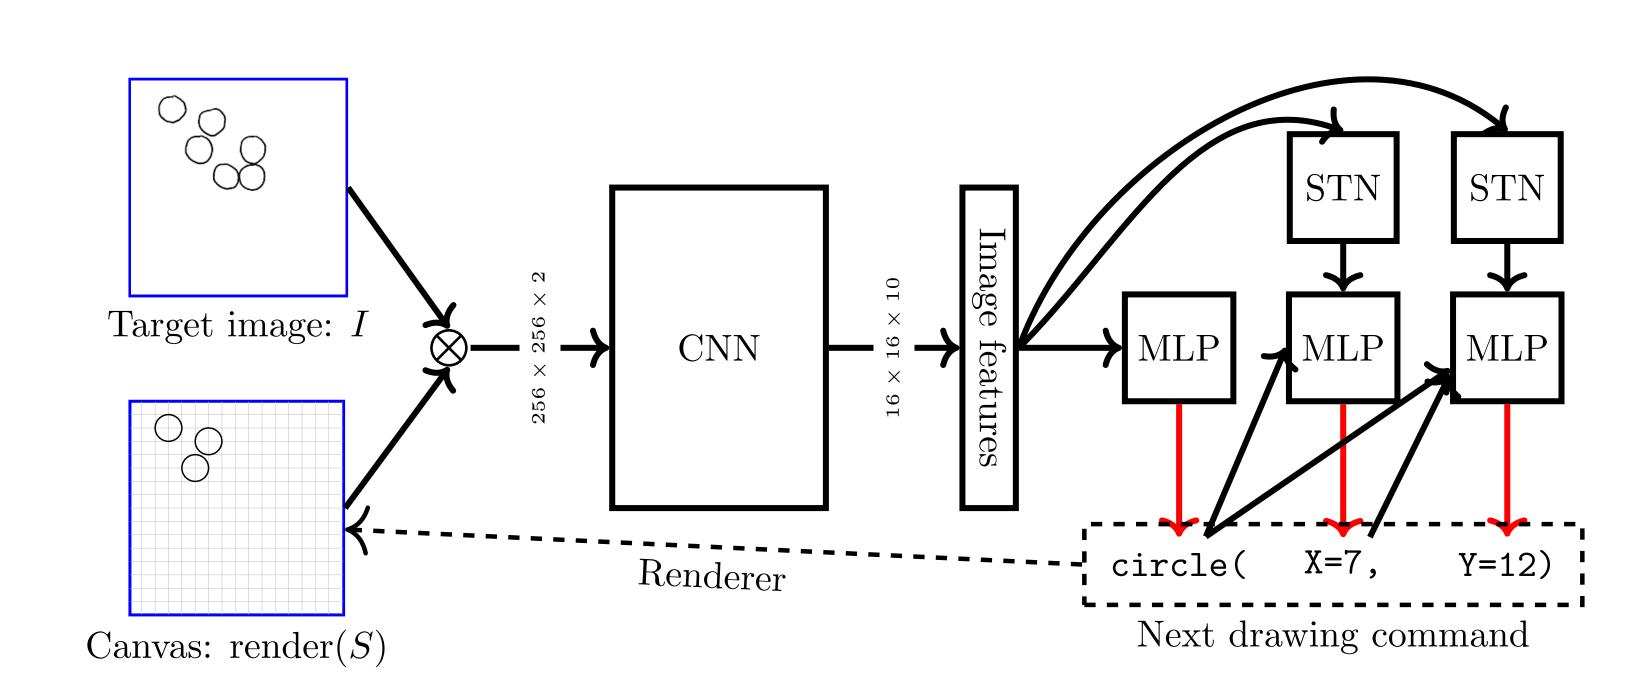
We are grateful for advice from Will Grathwohl and Jiajun Wu on the neural architecture. Funding from NSF GRFP, NSF Award #1753684, the MUSE program (DARPA grant FA8750-14-2-0242), and AFOSR award FA9550-16-1-0012.

Two-Stage Pipeline



Black arrows: Top-down generative model; Program→Spec→Image. Red arrows: Bottom-up inference procedure. Bold: Random variables (image/spec/program)

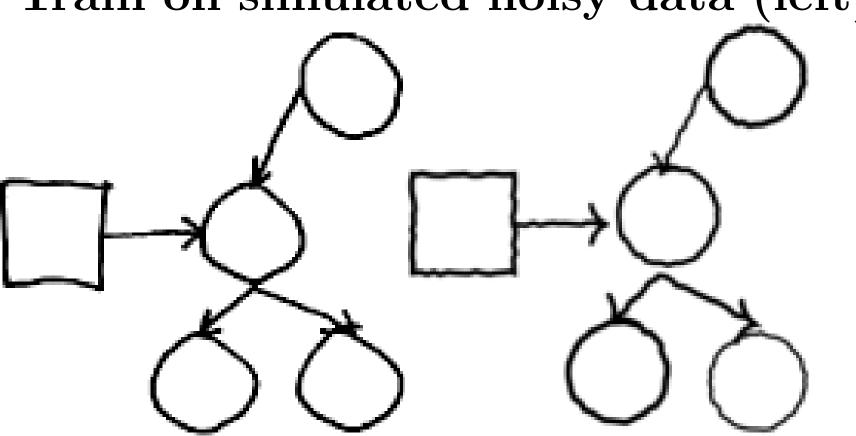
Parsing Images into Drawing Commands (specs)



Blue: network inputs. Black: network operations. Red: draws from a multinomial. Typewriter font: network outputs. Renders on a 16 × 16 grid, shown in gray. STN: differentiable attention mechanism. Combined with stochastic search (Sequential Monte Carlo)

Generalizing to real hand drawings

Train on simulated noisy data (left) & learn distance metric between images (right)

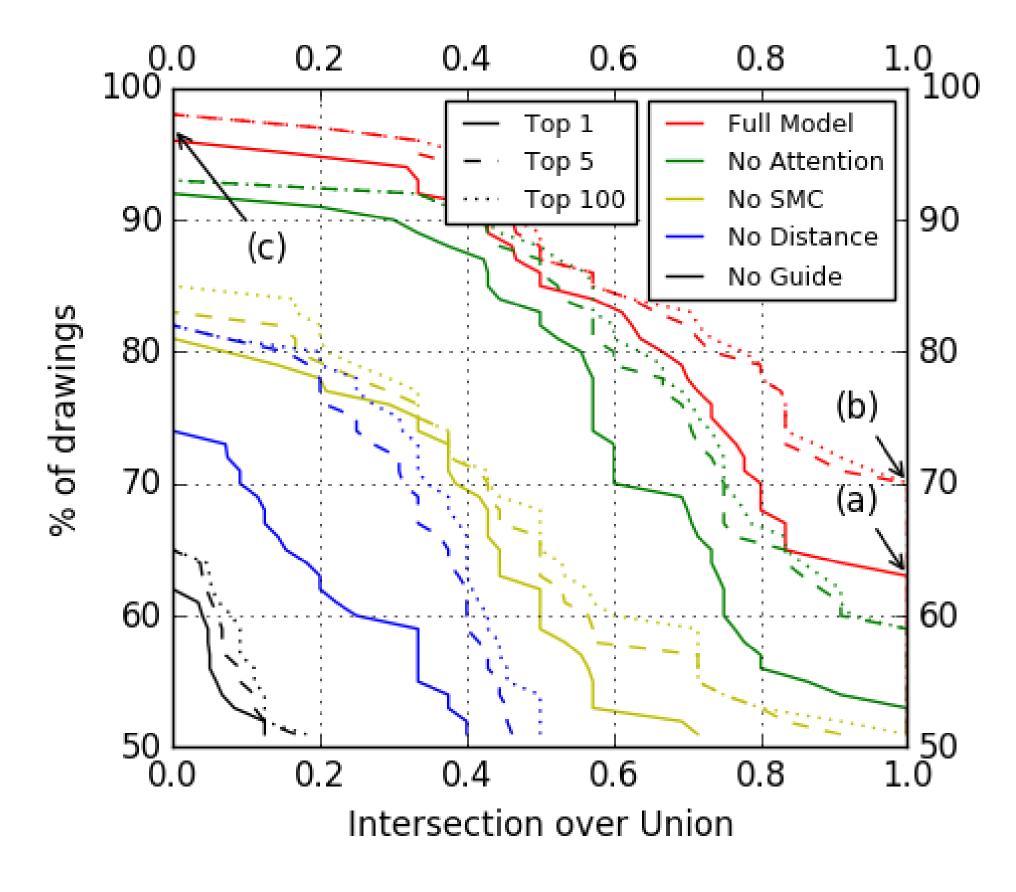


(a): hand drawing (b): noisy render of (a)

Learned distance metric

Serves as likelihood surrogate for SMC $-\log L_{\mathrm{learned}}(\mathrm{render}(S_1)|\mathrm{render}(S_2)) \approx |S_1 - S_2| + |S_2 - S_1|$ S_1 : noisy render of random scene

 S_2 : nonnoisy render of random scene



(Left) NN outputs vs ground truth on hand drawings (measured by IoU), as we consider larger sets of samples (1, 5, 100). (a) for 63% of drawings the model's top prediction is exactly correct; (b) for 70% of drawings the ground truth is in the top 5 model predictions; (c) for 4% of drawings all of the model outputs have no overlap with the ground truth. Red: the full model. Other colors: ablations. Model is at ceiling for synthetic ETEX output.

Domain-Specific Language for Graphics Programs

We	allow	loops	(for)	with	conditionals	$(\mathtt{if}),$	vertical/horizo	ontal reflec-		
tions	(ref	lect),	variables	(Var)	and	affine	transformations	$(\mathbb{Z} \times \text{Var} + \mathbb{Z}).$		
Prog	$\operatorname{gram} \to \mathfrak{S}$	Statement:	···; Stater	nent						
State	$ment \rightarrow c$	$\mathtt{circle}(\mathrm{Ex}$	xpression,Ex	(xpression						
State	$Statement \rightarrow rectangle(Expression, Expression, Expression, Expression)$									
$Statement \rightarrow line(Expression, Expression, Expression, Expression, Boolean, Boolean)$										
State	$ment \rightarrow z$	$for(0 \le V)$	ar < Expres	ssion) {	if $(Var > 0)$) { Progr	ram }; Program }	>		
State	$\mathrm{ment} \rightarrow 1$	reflect(A	Axis) { Pro	ogram }						
Expre	ssion \rightarrow 2	$\mathbb{Z} \times \text{Var+} \mathbb{Z}$								
	$Axis \rightarrow X$	$X = \mathbb{Z} \mid Y$	$= \mathbb{Z}$							
	$\mathbb{Z} \to \epsilon$	an integer								

Learning & Constraint-Based Program Synthesis

Sketch: state-of-the-art program synthesizer. Solar-Lezama 2008. Solves, for spec S & program p:

$$\operatorname{program}(S) = \underset{p \in \mathrm{DSL, s.t. } p \text{ consistent w/} S}{\operatorname{arg min}} \operatorname{cost}(p)$$

Learn policy π_{θ} to accelerate program synthesizer's search. $\pi_{\theta}(\cdot|S) \in \Delta^{\Sigma}$, where $\Sigma \ni \sigma$ a set of synthesis problem (i.e., $\sigma \in \Sigma$ is a sketch) Inference strategy: Timeshare according to $\pi_{\theta}(\cdot|S)$, like in Levin Search

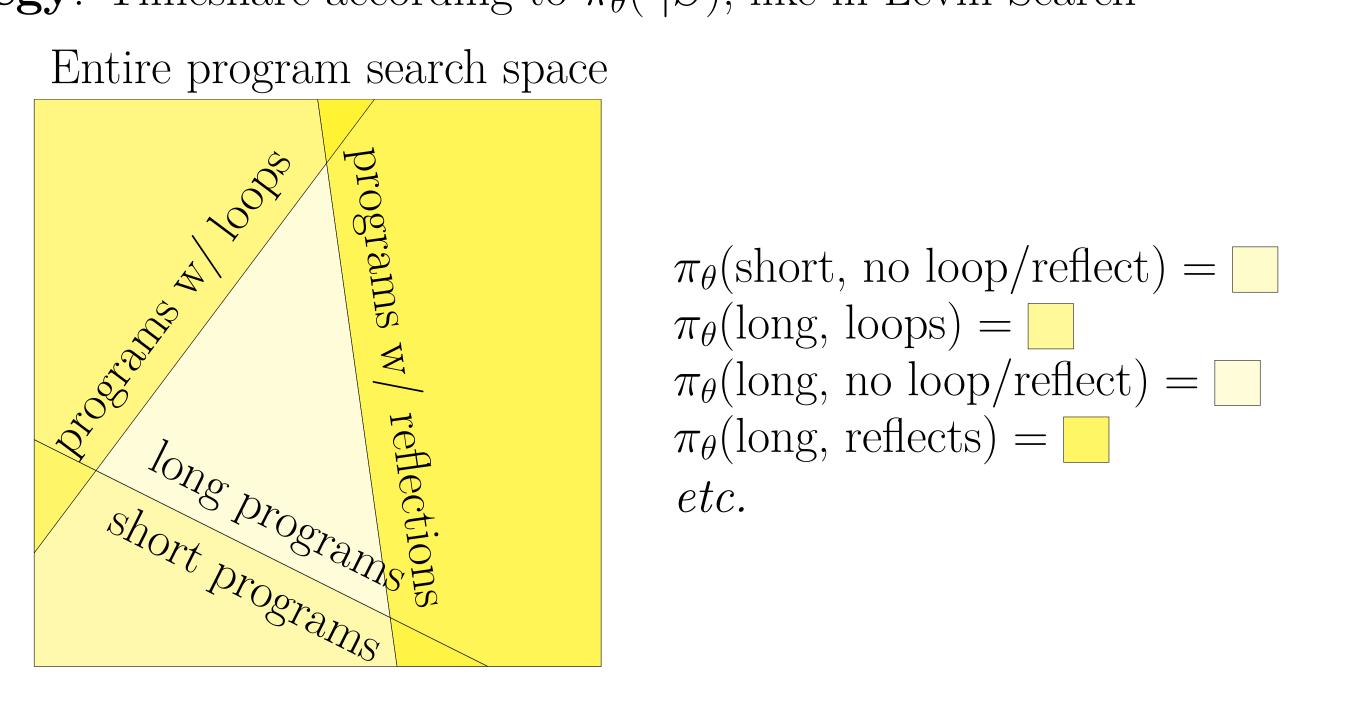


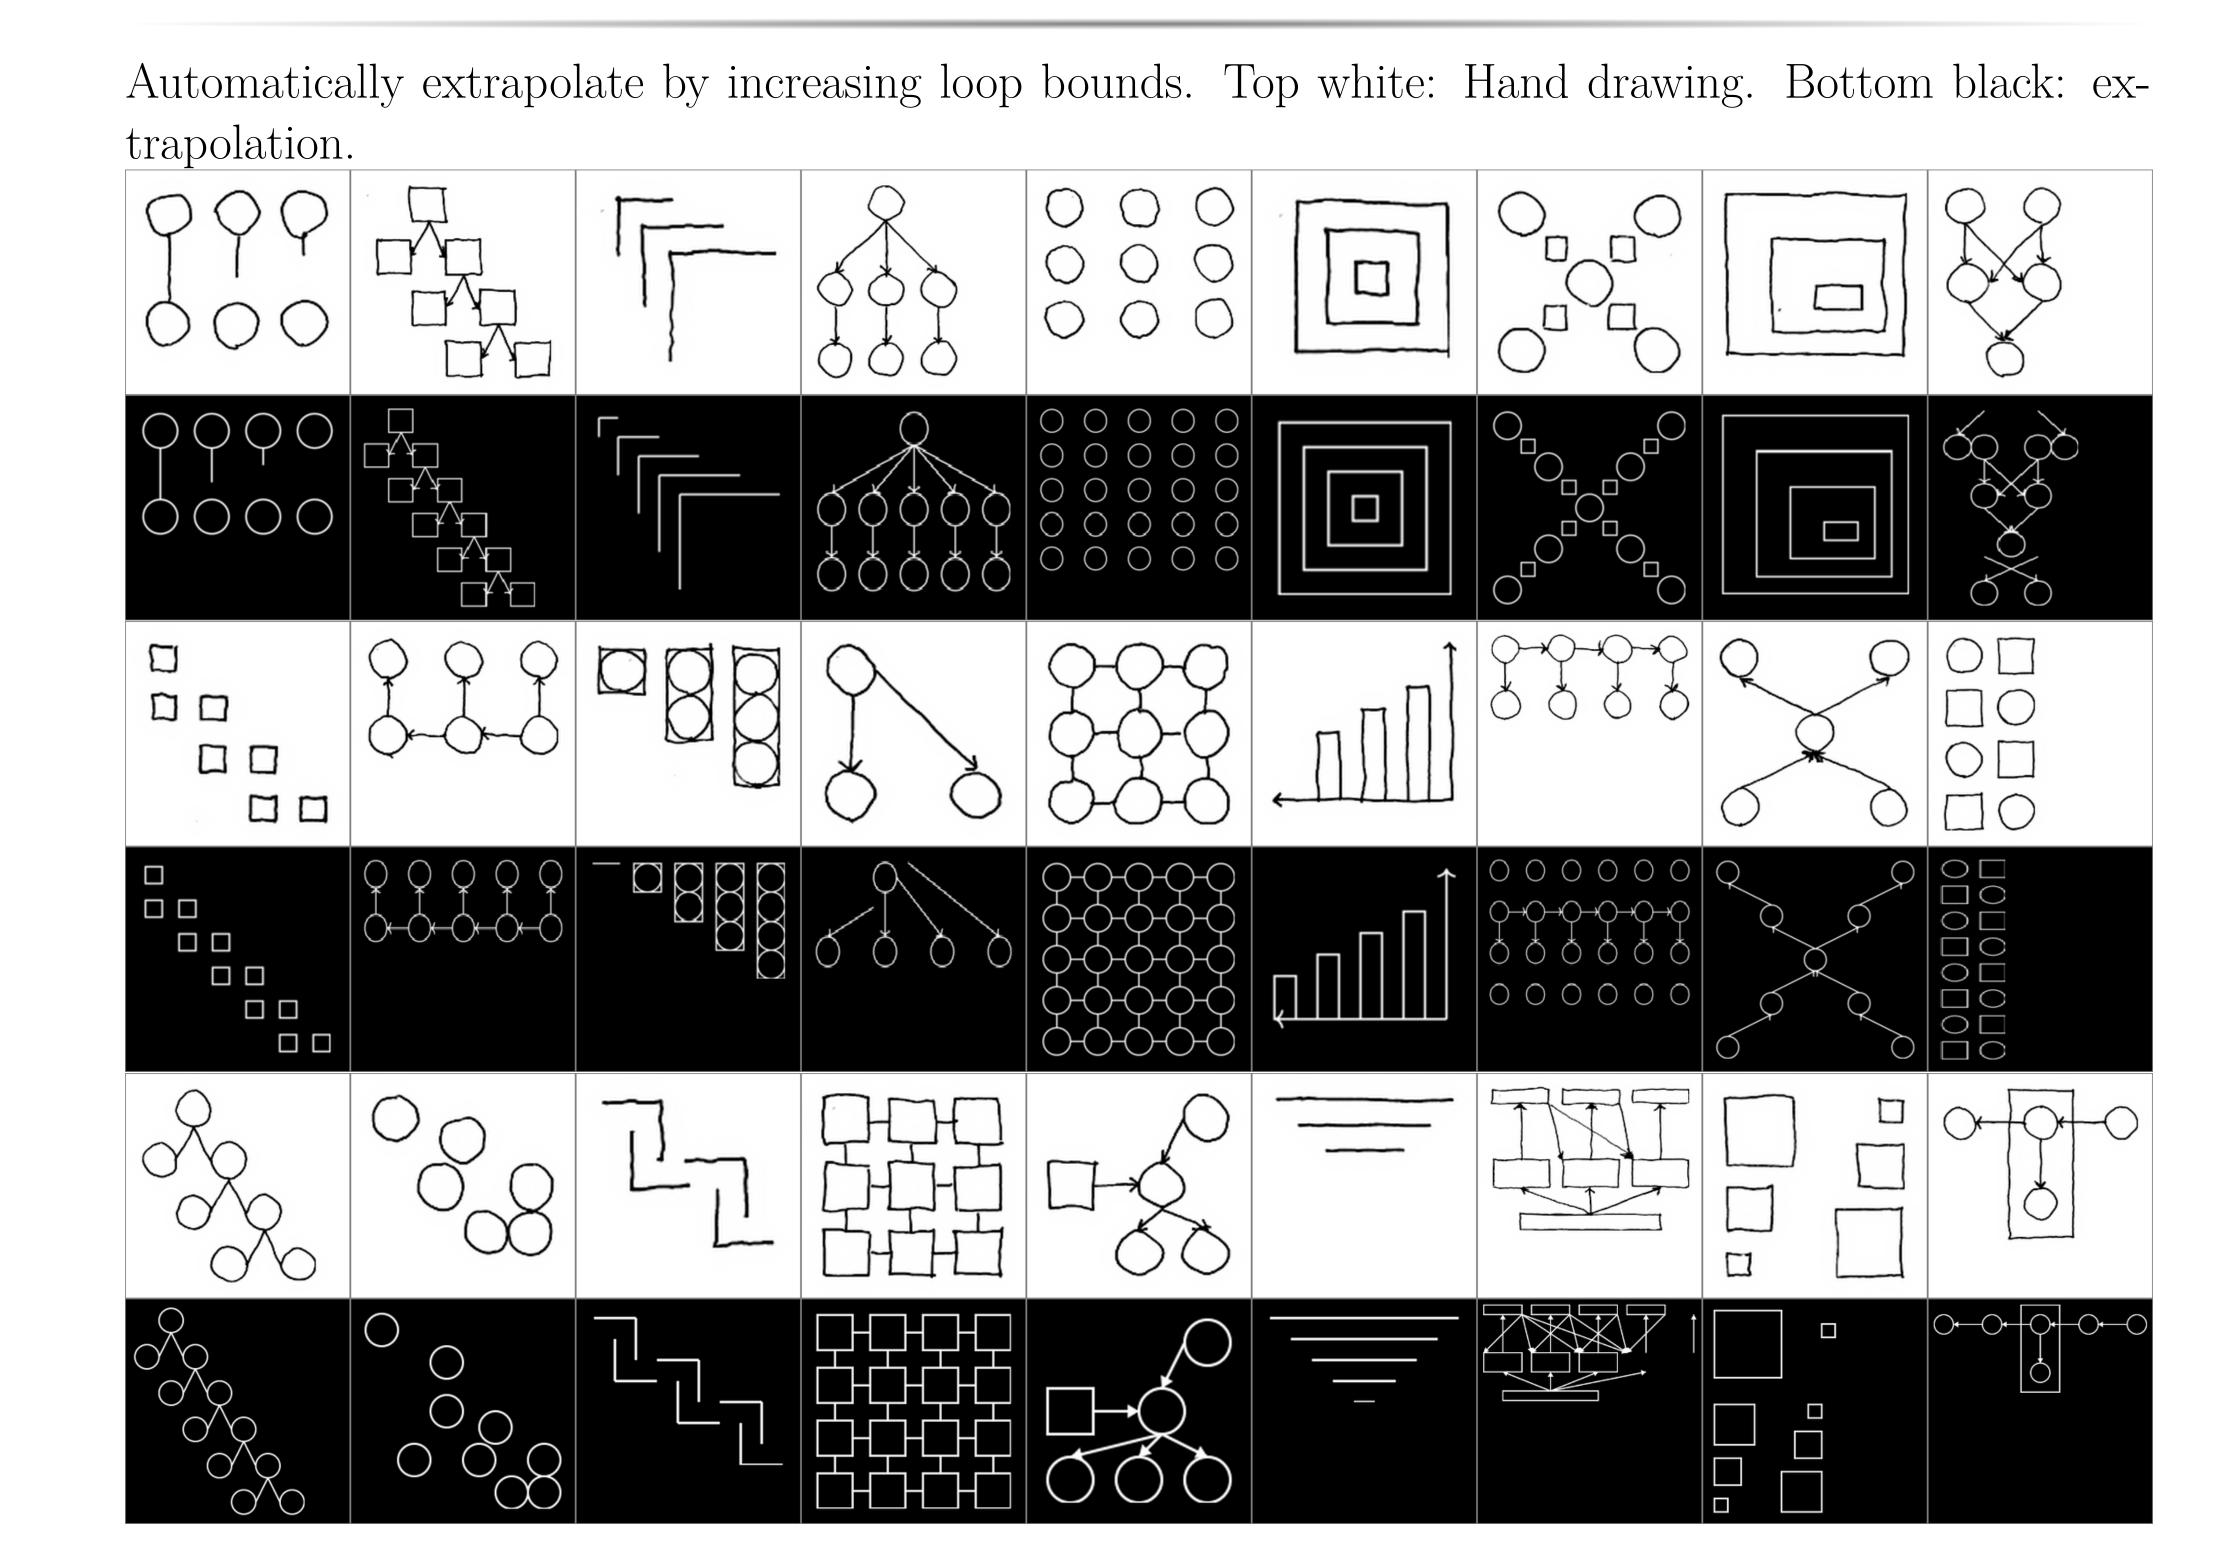
Figure 1: The bias-optimal search algorithm divides the entire (intractable) program search space in to (tractable) program subspaces (written σ), each of which contains a restricted set of programs. For example, one subspace might be short programs which don't loop. The policy π predicts a distribution over program subspaces: weight assigned by π indicated by shading

Differentiable loss (\mathcal{D} a corpus of synthesis problems):

$$Loss(\theta; \mathcal{D}) = \mathbb{E}_{S \sim \mathcal{D}} \left[\min_{\sigma \in BEST(S)} \frac{t(\sigma|S)}{\pi_{\theta}(\sigma|S)} \right] + \lambda \|\theta\|_{2}^{2}$$
 where $\sigma \in BEST(S)$ if a minimum cost program for S is in σ . (1)

Experimental results on synthesis policy: See lower right of poster

Extrapolating Drawings



Error correction

'Top down' influences upon perception: reasoning engine (program synthesizer) can influence agent's percept through higher-level considerations like symmetry and alignment.

 $\hat{S}(I) = \arg\max L_{\text{learned}}(I|\text{render}(S)) \times \mathbb{P}_{\theta}[S|I] \times \mathbb{P}_{\beta}[\text{program}(S)]$

$$\beta^* = \arg\max_{\beta} \mathbb{E} \left[\log \frac{\mathbb{P}_{\beta}[\operatorname{program}(S)] \times L_{\operatorname{learned}}(I|\operatorname{render}(S)) \times \mathbb{P}_{\theta}[S|I]}{\sum_{S' \in \mathcal{F}(I)} \mathbb{P}_{\beta}[\operatorname{program}(S')] \times L_{\operatorname{learned}}(I|\operatorname{render}(S')) \times \mathbb{P}_{\theta}[S'|I]} \right]$$

where $\mathcal{F}(I)$ is set of parses output by neural net on image I; $\mathbb{P}_{\beta}[\cdot]$ is prior over programs parameterized by β ; and expectation is taken over a corpus of program synthesis problems.

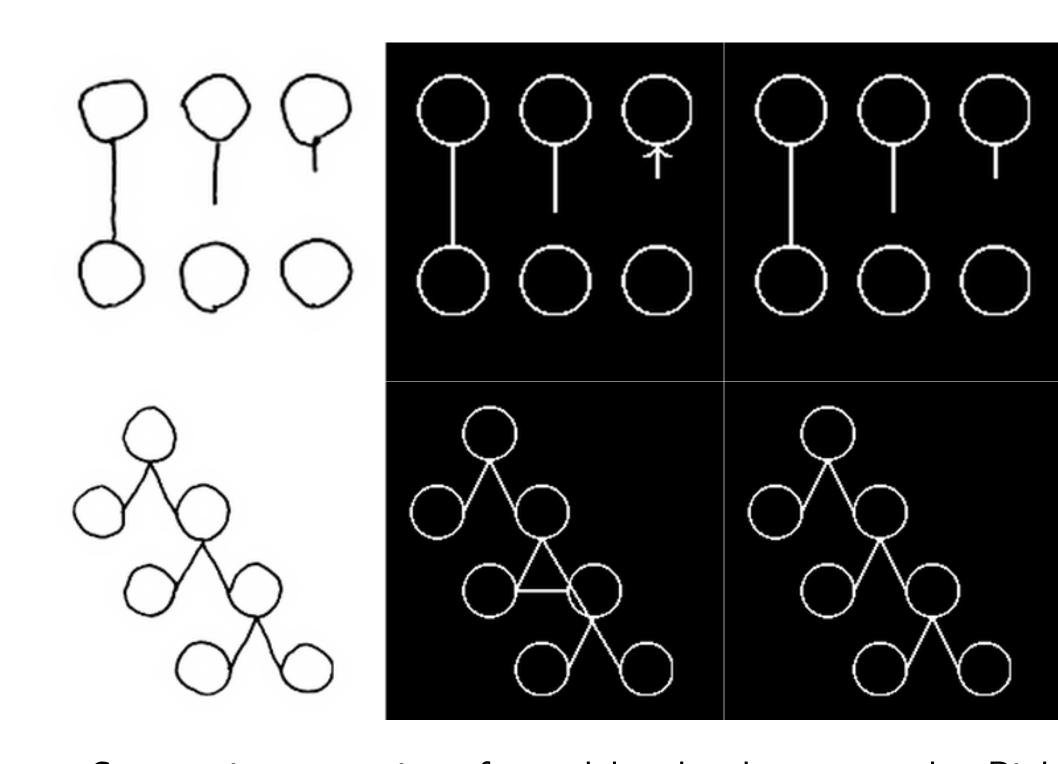
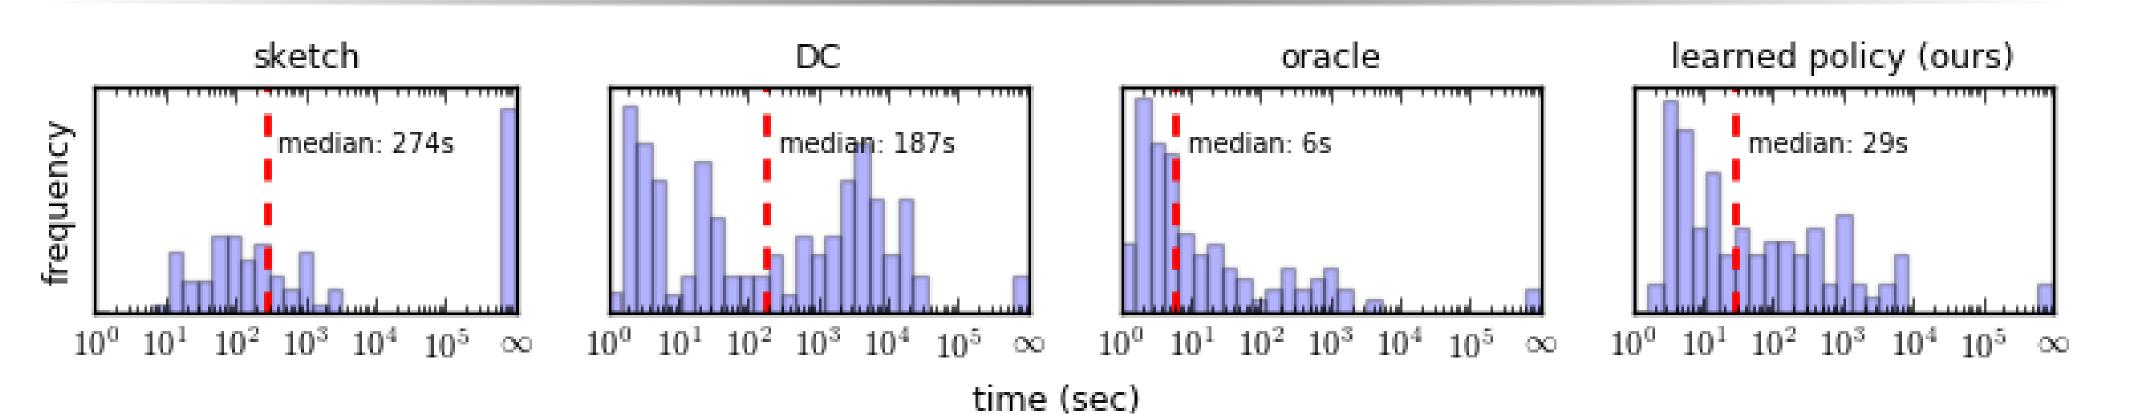


Figure 2: Left: hand drawings. Center: interpretations favored by the deep network. Right: interpretations favored after learning a prior over programs. The prior favors simpler programs, thus (top) continuing the pattern of not having an arrow is preferred, or (bottom) continuing the "binary search tree" is preferred.

Synthesis times



Time to synthesize a minimum cost program. Sketch: out-of-the-box performance of Sketch. DC: Deep-Coder style baseline that predicts program components, trained like Balog 2016. Oracle: upper bounds the performance of any bias-optimal search policy. ∞ = timeout. Red dashed line is median time