
Supplement to: Inferring Graphics Programs from Images

Anonymous Author(s)

Affiliation

Address

email

1 Simulating hand drawings

We introduce noise into the rendering process by:

- Rescaling the image intensity by a factor chosen uniformly at random from $[0.5, 1.5]$
- Translating the image by ± 3 pixels chosen uniformly random
- Rendering the \LaTeX using the `pencildraw` style, which adds random perturbations to the paths drawn by \LaTeX in a way designed to resemble a pencil.
- Randomly perturbing the positions and sizes of primitive \LaTeX drawing commands

2 Neural networks for guiding SMC

Let $L(\cdot|\cdot) : \text{image}^2 \rightarrow \mathcal{R}$ be our likelihood function: it takes two images, an observed target image and a hypothesized program output, and gives the likelihood of the observed image conditioned on the program output. We want to sample from:

$$[p|x] \propto L(x|\text{render}(p))[p] \quad (1)$$

where $[p]$ is the prior probability of program p , and x is the observed image.

Let p be a program with L lines, which we will write as $p = (p_1, p_2, \dots, p_L)$. Assume the prior factors into:

$$[p] \propto \prod_{l \leq L} [p_l] \quad (2)$$

Define the distribution $q_L(\cdot)$, which happens to be proportional to the above posterior:

$$q_L(p_1, p_2, \dots, p_{L-1}, p_L) \propto q_{L-1}(p_1, p_2, \dots, p_{L-1}) \times \frac{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}, p_L))}{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}))} \times [p_L] \quad (3)$$

Now suppose we have some samples from $q_{L-1}(\cdot)$, and that we then sample a p_L from a distribution proportional to $\frac{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}, p_L))}{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}))} \times [p_L]$. The resulting programs p are distributed according to q_L , and so are also distributed according to $[p|x]$.

How do we sample p_L from a distribution proportional to $\frac{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}, p_L))}{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}))} \times [p_L]$? We have a neural network that takes as input the target image x and the program so far, and produces a distribution over next lines of code (p_L). We write $\text{NN}(p_L|p_1, \dots, p_{L-1}; x)$ for the distribution output by the neural network. So we can sample from NN and then weight the samples by:

$$w(p_L) = \frac{[p_L]}{\text{NN}(p_L|p_1, \dots, p_{L-1}; x)} \times \frac{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}, p_L))}{L(x|\text{render}(p_1, p_2, \dots, p_{L-1}))} \quad (4)$$

Algorithm 1 Neurally guided SMC

Input: Neural network NN, beam size N , maximum length L , target image x

Output: Samples of the program trace

Set $B_0 = \{\text{empty program}\}$

for $1 \leq l \leq L$ **do**

for $1 \leq n \leq N$ **do**

$p_n \sim \text{Uniform}(B_{l-1})$

$p'_n \sim \text{NN}(\text{render}(p), x)$

 Define $r_n = p'_n \cdot p_n$

 Set $\tilde{w}(r_n) = \frac{L(x|r_n)}{L(x|p_n)} \times \frac{[p'_n]}{[p'_n = \text{NN}(\text{render}(p), x)]}$

end for

 Define $w(p) = \frac{\tilde{w}(p)}{\sum_{p'} \tilde{w}(p')}$

 Set B_l to be N samples from r_n distributed according to $w(\cdot)$

end for

return $\{p : p \in B_{l \leq L}, p \text{ is finished}\}$

23 Then we can resample from these now weighted samples to get a new population of particles (here
24 programs are particles), where each program now has L lines instead of $L - 1$.

25 This procedure can be seen as a particle filter, where each successive latent variable is another line of
26 code, and the emission probabilities are successive ratios of likelihoods under $L(\cdot|\cdot)$.

27 **Comments for Dan.** Right now I'm not actually sampling from the neural network - instead, I
28 enumerate the top few hundred lines of code suggested by the network, and then weight them by their
29 likelihoods. So actually the form of NN is:

$$\text{NN}(p_L | p_1, \dots, p_{L-1}; x) \propto \begin{cases} 1, & \text{if } p_L \in \text{top hundred neural network proposals} \\ 0, & \text{otherwise.} \end{cases} \quad (5)$$

30 Do you think this is a problem? The neural network puts almost all of its mass on a few guesses. In
31 order to get the correct line of code I sometimes need to get something like the 50th top guess, so I
32 don't want to literally just sample from the distribution suggested by the neural network.